Causation, Meaning, and Communication

Ari Beller

Department of Psychology, Stanford University

Tobias Gerstenberg

Department of Psychology, Stanford University

September 20, 2024

Author Note

Abstract

The words we use to describe what happened shape what comes to a listener's mind. How do speakers choose what causal expressions to use? How does that choice impact what listeners imagine? In this paper, we develop a computational model of how people use the causal expressions "caused", "enabled", "affected", and "made no difference". The model first builds a causal representation of what happened. By running counterfactual simulations, the model computes several causal aspects that capture the different ways in which a candidate cause made a difference to the outcome. Logical combinations of these aspects define a semantics for the causal expressions. The model then uses pragmatic inference to decide what word to use in context. We test our model in a series of experiments and compare it to prior psychological accounts. In a set of psycholinguistic studies, we verify the model's semantics and pragmatics. We show that the causal expressions exist on a hierarchy of specificity, and that participants draw informative pragmatic inferences in line with this scale. In the next two studies, we demonstrate that our model quantitatively fits participant behavior in a speaker task and a listener task involving dynamic physical scenarios. We compare our model to two lesioned alternatives, one which removes pragmatic inference, and another which removes semantics and pragmatics. Our full model better accounts for participants' behavior than both alternatives. Taken together, these results suggest a new way forward for modeling the relationship between language and thought in the study of causality.

Keywords: causality; counterfactuals; mental simulation; intuitive physics; language; semantics; pragmatics.

Causation, Meaning, and Communication

Introduction

The words we use matter. When you hear someone say "Tom killed Bill", an image of a murderous scene pops into your mind. If you hear instead that "Tom caused Bill to die", you might imagine a different scenario, one that's less direct and murderous than the first. Both of these sentences attribute a causal role to Tom in Bill's death, but the subtle differences in phrasing amount to worlds of difference in meaning. In everyday communication, causation enters into people's language in all kinds of innocuous yet impactful ways. While it is easy to miss the significance of these linguistic choices in the moment, their variety and flexibility supports people's capacity to effortlessly convey and comprehend causal stories that are both complex and specific.

Studying the language of causation is a multi-disciplinary endeavor. Linguists aim to understand how people talk about cause and effect (Aronson, 1971; Garvey & Caramazza, 1974; Hobbs, 2005; Levin & Hovav, 1994; Shibatani, 1976; Talmy, 1988) and in doing so often draw inspiration from the philosophical literature (Dowe, 2000; Hall, 2004; Lewis, 1973). Psychologists study how people represent causal relationships and how they choose specific causal expressions on the basis of that knowledge (Cheng & Novick, 1991; Gerstenberg, Goodman, Lagnado, & Tenenbaum, 2021; Rose, Sievers, & Nichols, 2021; Wolff, 2007). In this paper we develop the *counterfactual simulation model of causal language* which incorporates insights from these three disciplines. Our model combines a psychological approach for representing causation that is grounded in the philosophical literature on causation, with linguistic tools for modeling meaning and pragmatic communication. The combination of these techniques helps us shed light on the interaction of language and thought in how people communicate about causality. Before describing how the model works, we discuss the relevant background literature from each of these three disciplines: philosophy, linguistics, and psychology.

The philosophy of causal language

There are two major philosophical approaches for analyzing causation: *dependence theories* and *process theories.* According to dependence theories, causality is a dependence relation between cause and effect. Dependence has been characterized in various ways. For example, in counterfactual theories, A is a cause of B if A and B happened, and it's true that B would not have happened if A had not happened (Lewis, 1973; Mackie, 1974; Woodward, 2003). Imagine that Marco threw a stone at a window and the window broke. Here, Marco's throwing the stone caused the window to break because both of these events happened, and because the window wouldn't have broken had Marco not thrown the stone. On the other hand, process theories claim that causal relationships are defined by spatiotemporally continuous processes that link causes with effects (Dowe, 2000; Machamer, Darden, & Craver, 2000; Salmon, 1984). So, A is a cause of B if A transferred some property, such as energy or force, to B. Here, Marco caused the window to break because the stone transferred force to the breaking window via a spatiotemporally contiguous process.

Philosophers are not only interested in metaphysical questions about what causation is. They also care about developing theories that accord with human causal thinking. Writing from the dependence tradition, Woodward (2021) argues that these two aspects of understanding causality are in fact deeply intertwined. According to him, it would be hard to make sense of the general success people have using causal reasoning in their daily lives if it didn't correspond in some meaningful ways to the actual causal structure of the world. In this way, causal psychology very likely tells us something about metaphysical causality. At the same time, he points out that metaphysical accounts of causation generally rely on human conceptions of what is and is not causal to greater or lesser extent. As Woodward notes, many philosophers when assessing whether their causal theory is adequate often rely on intuitions about what people would say in some causal scenario. If these intuitions about people's causal judgments match a theory's predictions this is taken as evidence in

support of the theory. This approach is common throughout the literature on causality (Hall, 2007; Halpern, 2016; Halpern & Pearl, 2005; Hitchcock, 2009; Hitchcock & Knobe, 2009), and though the references to human judgments may be more or less explicit, they underlie the approach and emphasize the importance of causal psychology for understanding causality more generally.

Other philosophers are more explicitly focused on understanding causation by analyzing people's causal language. Working in the process tradition, Fair (1979) states that a central mystery in understanding causation is accounting for the broad consistency in people's agreement to causal statements. According to (Fair, 1979, pg. 220), solving this mystery requires "stating explicit truth-conditions for simple declarative sentences containing the word, 'cause'". Fair doesn't offer these truth-conditions himself, but he suggests a common feature that people are sensitive to when they interpret situations as causal. According to Fair, causation is reducible to energy-momentum flow from cause to effect. In a game of pool, a cue-ball causes an eight-ball to sink because the cue-ball transferred energy to the eight-ball upon collision, and that energy carried the eight-ball to into the pocket (but see Hitchcock, 1995).

Philosophers not only analyze the meaning of "cause", they also care about the differences between various causal expressions such as "cause" versus "affect" (McDermott, 1995). The subtle differences between causal expressions have animated debates in moral philosophy. For example, philosophers differentiate "killing" from "letting die" on the basis of causal structure (Foot, 1967; Malm, 1989; McGrath, 2003; McMahan, 1993; Thomson, 1976b). The distinctions that drive this discussion rest on the causal status of omissive causes (causing something to happen by "not preventing" it from happening) and sufficient causes (Gerstenberg & Stephan, 2021; Mackie, 1974; McGrath, 2003; Schaffer, 2000). Whether we say "killing" or "letting die" has practical significance in bioethics, and the meanings of these words are wrapped up in broader debates about end-of-life care in medicine (Rodríguez-Arias, Rodriguez Lopez, Monasterio-Astobiza, & Hannikainen, 2020),

a woman's right to an abortion (Thomson, 1976a), and the moral obligation to vaccinate (Flanigan, 2014).

The linguistics of causal language

Linguists have studied the many ways in which causality reveals itself in language both implicitly (e.g. Garvey & Caramazza, 1974; Hartshorne, 2013; Niemi, Hartshorne, Gerstenberg, Stanley, & Young, 2020) and explicitly (e.g. Kaufmann, 2013; Talmy, 1988). Causative constructions come in different varieties. Two of the most common in English are periphrastic causatives, where the causal meaning is expressed in a domain-general verb (e.g., "Jane **caused** the ice to break"), and lexical causatives, where the causal meaning is embedded in a domain-specific verb (e.g., "Jane **broke** the ice"). The semantic relationships between lexical causatives and corresponding periphrastic causatives with analogous meanings ("killed" vs. "caused to die") is a topic of extensive discussion (Cruse, 1972; Fodor, 1970; Shibatani, 1976; Smith, 1970; Wierzbicka, 1975). Though at first glance, lexical causatives and corresponding periphrastic paraphrases might seem to mean the same thing, linguists have noted that their meanings can come apart. Katz (1970) provides an example where a sheriff is set to duel with an outlaw. Prior to the duel, the sheriff has his gun poorly repaired by a gunsmith, such that when the sheriff ultimately faces the outlaw, the gun doesn't fire and the sheriff is shot dead. According to Katz, the gunsmith caused the sheriff to die, but he didn't kill him. The "killing" description is reserved for the outlaw who actually shot the sheriff.

A common theme emerging from this literature is that lexical causatives imply some form of direct causation over and above the general causal relationships implied by periphrastic causatives. It is appropriate to say that the outlaw killed the sheriff (because the actions of the outlaw are the direct cause of the sheriff's death), while the gunsmith merely caused the sheriff to die. Baglini and Siegal (2021) model the semantic differences between lexical and periphrastic causatives using SEMs. Baglini and Siegal argue that the different causative constructions imply different structural roles of causes cited in an

underlying SEM. Periphrastic causatives (specifically the periphrastic causative "cause") can be applied to any variable in the SEM that is a necessary condition for some observed outcome. However, to qualify as the subject of a lexical causative, a variable in the SEM must be part of a set of causal conditions that are jointly sufficient for bringing about the outcome. Baglini and Siegal show that these definitions account for the intuitions of direct causation that have shaped many earlier accounts, while also addressing counterexamples to the directness criteria that have been raised more recently (see Wolff, 2003, for a review of different approaches to direct causation). In a similar vein, Nadathur and Lauer (2020) use the SEM framework to explain semantic differences between "caused" and "made", whereby "caused" is analyzed in terms of necessity, and "made" in terms of sufficiency.
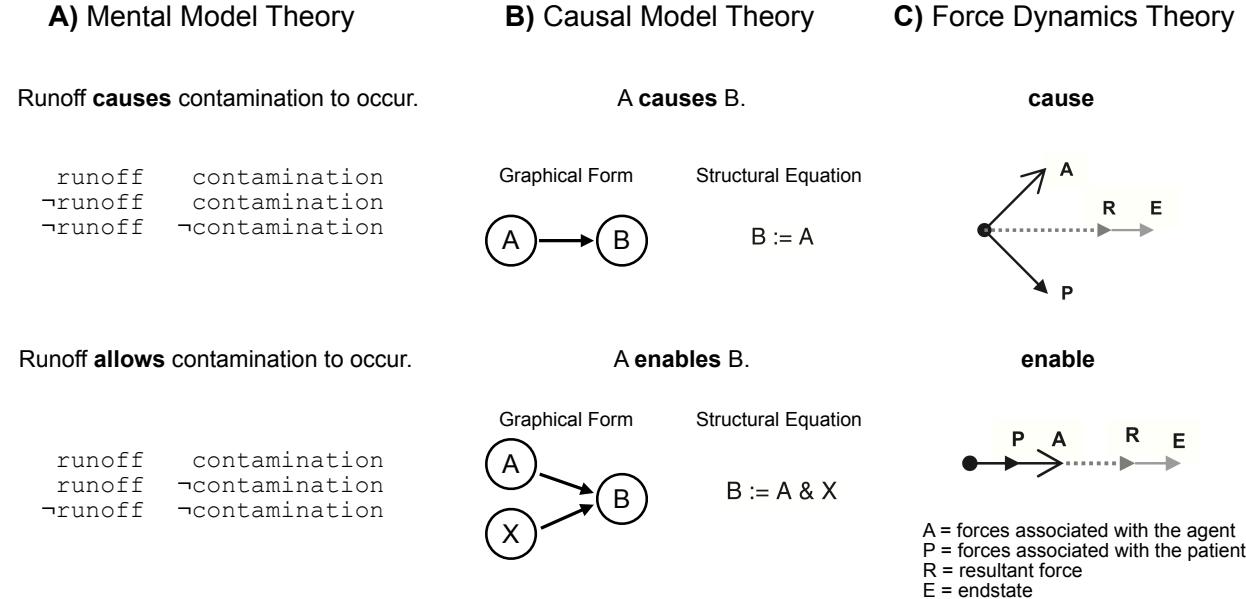
The SEM approaches to causal semantics are rooted in dependence theories of causation. Within the process theory tradition, Talmy (1988) developed the force-dynamics model for analyzing the meaning of various periphrastic causatives. A force-dynamic description identifies a focal entity called an agonist and an opposing entity called an antagonist along with their intrinsic tendencies toward action or rest, their relative strengths, and a resultant action. For example, in a situation where a ball knocks into a table lamp and the lamp falls over, the agonist would be the table lamp which has a tendency toward rest, while the antagonist would be the ball which has a tendency toward action. In this case, the ball's tendency toward action overpowers the lamp's tendency toward rest leading to a resultant action where the lamp falls over. According to Talmy, this particular force-dynamic configuration corresponds to the prototypical causation scenario – a scenario that people would describe using the periphrastic causative "cause".

The work discussed so far focuses on the semantics of causative constructions. Pragmatic inferences, however, also impact their interpretation (Degen, 2023; Goodman & Frank, 2016; Grice, 1975; Schaffer, 2013). McCawley (1978) suggests that principles of cooperative conversation influence people's selection of causative constructions in context. For example, the presence of alternatives affects people's inferences about the directness of

causation. Upon hearing that "the gunsmith caused the sheriff to die", a listener is likely to infer that the gunsmith's causal role was indirect because the speaker could have used the alternative "killed" to communicate a more direct causal role. However, when there is no lexical alternative, such as when "Bill caused Mary to laugh", the periphrastic causative is acceptable regardless of whether causation was direct or indirect. McCawley argues that inferring indirect causation from periphrastic causatives isn't due to their semantics, but rather that it's a type of pragmatic inference called a conversational implicature (Grice, 1975). Listeners in conversation generally assume that speakers are as informative as they can be, and if they choose a less specific description ("caused to die") when a more specific one is available ("killed"), this suggests that the more specific situation isn't true.

The psychology of causal language

Psychologists too have shown interest in people's use of causal language, paying particular attention to the mental representations that underlie people's use of different causal expressions. In this space, a number of researchers have adopted a modeling approach where they take a set of causal expressions (usually periphrastic causatives) and a modeling language and show that people's use of a particular expression corresponds to a particular mental representation expressed in the modeling language. An example of this approach is Goldvarg and Johnson-Laird (2001), who use mental model theory (Johnson-Laird, 1989) to explain people's use of the periphrastic causatives "cause" and "allow" (see also Khemlani et al., 2014). Mental model theory construes mental representations in terms of logical possibilities. Different periphrastic causatives can be distinguished by the sets of possibilities that they permit and preclude. Figure 1a illustrates the mental model representations for "cause" and "allow". The sentence "Runoff causes contamination to occur" permits three possibilities: both runoff and contamination occur, no runoff occurs (indicated by the negation sign) but contamination still occurs, or neither runoff nor contamination occur. The possibility that is precluded is that runoff occurs but contamination does not occur. According to this definition of "cause", whenever

A) Mental Model Theory   **B)** Causal Model Theory   **C)** Force Dynamics Theory

Runoff **causes** contamination to occur.   A **causes** B.   **cause**

```
 runoff     contamination
¬runoff     contamination
¬runoff    ¬contamination
```

Graphical Form   Structural Equation

A ⟶ B   B := A

Runoff **allows** contamination to occur.   A **enables** B.   **enable**

Graphical Form   Structural Equation

```
 runoff     contamination
 runoff    ¬contamination
¬runoff    ¬contamination
```

A
       B   B := A & X
X

A = forces associated with the agent
P = forces associated with the patient
R = resultant force
E = endstate

Figure 1. Prior psychological models distinguish among the meanings of different causal verbs by articulating different mental representations that underlie a particular verb's use. Different accounts use different modeling approaches to describe the mental representations. Goldvarg and Johnson-Laird (2001) and Khemlani, Barbey, and Johnson-Laird (2014) use Mental Model Theory, a logical representation that specifies the combinations of cause and effect that are consistent with a particular causal verb. Sloman, Barbey, and Hotaling (2009) use Causal Models, a graphical formalism for describing how causes determine effects. Wolff (2007) uses Force Dynamics Theory, a process oriented approach that distinguishes between different causal concepts with different force configurations represented as vector diagrams. All three models provide definitions for the verb "caused". Mental Model Theory provides a definition for the verb "allow" while Causal Model Theory and Force Dynamics Theory provide a definition for the related verb "enabled".

the cause occurs the effect must follow. However, the effect can still occur in the absence of the cause. In other words, this definition implies that the cause is sufficient but not necessary for the outcome.

According to mental model theory, if a person's representation of the relationship between runoff and contamination corresponds to this set of possibilities and impossibilities, they will endorse the claim that runoff causes contamination to occur. On the other hand, if they endorse the sentence "Runoff allows contamination to occur", it suggests they think a different set of possibilities captures the relationship between runoff

and contamination. Like "cause", "allow" is consistent with the possibilities that runoff and contamination both occur, or both do not occur. However, unlike "cause", it permits the possibility that runoff occurs and contamination does not occur, and it precludes the possibility that runoff does not occur and contamination occurs. According to this definition, the effect cannot take place without the allowing event, however the presence of the allowing event does not imply that effect takes place. In other words, this definition of "allow" implies that the allowing event is necessary but not sufficient for the outcome.

Another approach for studying causal language was developed by Sloman et al. (2009) and uses causal models (Pearl, 2000; Sloman, 2005). Causal models are a graphical formalism that represents the causal relationships among variables (nodes) as directed edges from a cause variable to an effect variable. The values of the causal variables determine the values of the effect variables through functional relationships defined with structural equations. Sloman et al. (2009) leverage this modeling language to unpack the differences in mental representations that underlie people's use of the periphrastic causatives "cause" and "enable". Sloman et al. (2009) distinguish these mental representations as causal models with different structures.

Figure 1b illustrates Sloman et al.'s (2009) causal models for "cause" and "enable". According to their theory, the statement "A causes B" posits the simple causal relation with a single direct connection between the A and B variables. On the other hand, saying that "A enables B" implies a direct connection between A and B, as well as the existence of an additional accessory variable, X, which is also connected to B. The functional relationship that determines B's value in this case is conjunctive (indicated by the structural equation). B will be active in the event that both A and X are active.

Both Goldvarg and Johnson-Laird's (2001) mental model theory and Sloman et al.'s (2009) causal model theory sit within the dependence tradition of causality. They analyze the mental representations corresponding to particular causal expressions with models that represent relations of logical and structural dependence. In the tradition of process theories

of causation, Wolff (2007) developed an account that builds on Talmy's force dynamics theory (Talmy, 1988). Wolff analyzes the mental representations underlying people's use of the periphrastic causatives "cause" and "enable" in terms of different force dynamic configurations. A force dynamic configuration assumes an agent and a patient, each with their own force vector, as well as an endstate. The way in which force vectors of the agent and patient combine in relation to the endstate determines whether the situation is better described using "cause" or "enable".

Figure 1c shows the vector diagrams representing the force dynamic configurations for these two expressions. In the "cause" situation, the patient vector (P) is initially oriented away from the endstate (E), but when it combines with the agent vector (A), the resultant (R) is oriented toward the endstate. This shift in the patient's tendency away from the endstate and then toward the endstate defines the "cause" configuration. In the "enable" situation, the patient vector is initially oriented toward the endstate. The combination with the agent vector yields a resultant intensified in the direction of the endstate. In enabling situations, the patient is continually oriented toward the endstate and the agent merely helps the patient along.
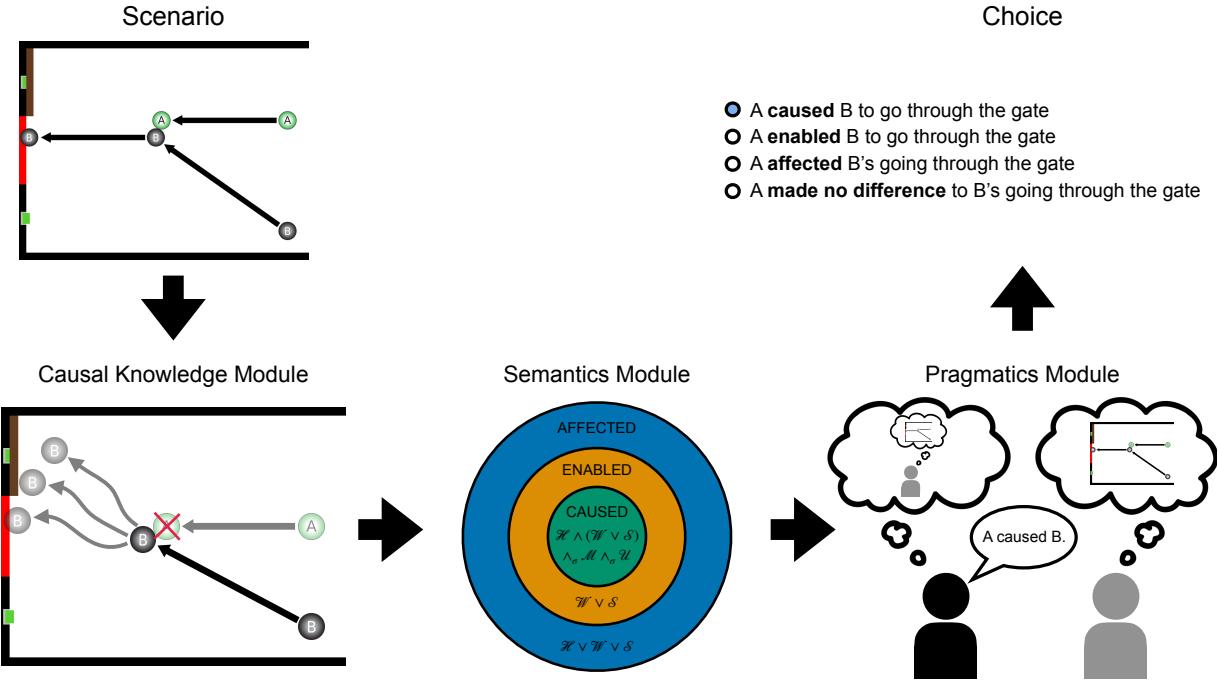
Mental models, causal models, and force dynamics theory have each shown success at capturing participants' judgments in their respective domains. However, all three approaches suffer from certain common shortcomings. As far as predicting participants' behavior, these models only provide a single modal prediction for the most likely utterance that a participant will select in a given situation. In spite of the fact that the experiments where these researchers tested their models all demonstrated variation in participants' responses on individual trials, the models can only predict a single response leaving these approaches with no way to capture the distribution of participant responses.

A further conceptual limitation is the way these models simplify the relationship between causal mental representations and linguistic processing. Each of these models associate particular periphrastic causatives directly with given mental representations

without considering how communicative goals and context shape the words people choose (Hilton, 1990). As we will demonstrate, this simplification can cloud semantic overlap in the causal expressions, and explicitly modeling this relationship between causal cognition and linguistic processing helps uncover this psychological subtlety.

Our primary goal in this paper is to develop a model that builds on the approach of this prior work, while addressing the limitations we have identified. Like previous models, our model aims to elucidate the relation between particular causal expressions and underlying mental representations. Extending these accounts, our model explicitly characterizes the relationship between mental representation and linguistic processing and provides quantitative predictions for each available utterance, allowing us to capture observed the distribution of participant responses. The combination of these developments provides a more nuanced depiction of the relationship between language and thought in the use of causal language, and uncovers subtleties in the meanings of causal expressions that have been hidden thus far.

The rest of the paper proceeds as follows. We begin by introducing our model and elaborating on how we link up psychological representations with tools from semantics and pragmatics to predict causal language use. We then present a series of experiments. The first is a psycholinguistics task where we validate our model's semantic and pragmatic assumptions. We demonstrate that the semantics of the causal expressions "caused" and "enabled" overlap in a way that aligns with our model semantics and contrasts with the assumptions of prior literature. Then in the second and third experiments, we test the model's quantitative predictions. We present participants with a speaker task (Experiment 2) and a listener task (Experiment 3). We compare our model's behavior to two alternative models that lesion different portions of the model capabilities. Our analysis reveals that the full model which includes causal knowledge, semantic representations, and pragmatic linguistic inference does the best job of explaining participant behavior. We close in the General Discussion by considering implications of the work and suggesting
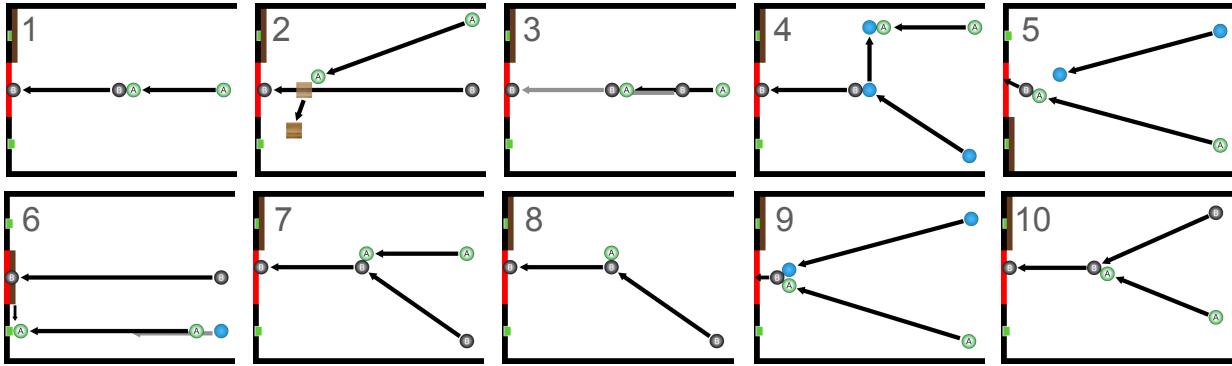
Figure 2. A schematic of the counterfactual simulation model in the speaker task. The model takes a scenario and runs different counterfactual simulations (represented by the transparent squiggly lines) to compute several aspects of causation that capture whether and how a candidate cause made a difference to the outcome (causal knowledge module). The meaning of various causal expressions is defined through logical combinations of these aspects of causation (semantics module). The model considers both what's true and what's informative when deciding what expression to choose (pragmatics module).

directions for future research.

A counterfactual simulation model of causal language

Our model combines causal reasoning with pragmatic communication to produce and understand causal language about particular events. The model has three components: a causality module, a semantics module, and a pragmatics module. Figure 2 provides an overview of how the model works. The causal knowledge module computes a causal representation of a scenario. This causal representation then feeds into the semantics module which determines which expressions are true in the scenario. Finally, on the basis of this semantics, a pragmatics module chooses an expression which is both true and informative for the given scenario. We discuss each of these components in turn and
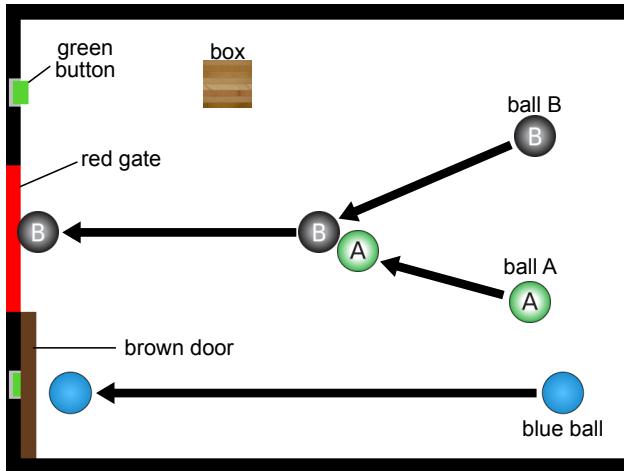
Figure 3. Sample scenarios from Experiments 2 and 3. 1) Classic Michottean case. Ball B sits in the middle of the scene and ball A comes in from the side and knocks it through the gate. 2) Ball B is headed toward the gate, and ball A knocks a box blocking its path out of the way. 3) Ball B is moving toward the gate unobstructed. Ball A comes up from behind and pushes it along, speeding it up. 4) The blue ball knocks ball B through the gate. Afterward the blue ball collides with ball A. 5) A case of causal preemption. Ball A knocks ball B through the gate shortly before the blue ball would have done the same. 6) Similar to scenario 5 but without direct contact. Here ball A pushes the button opening the gate shortly before the blue ball would have done the same. Opening the gate allows ball B to pass through the gate. 7) Ball A and ball B enter from the right side. The collision redirects ball B through the gate. 8) Similar to scenario 7 except here ball A is stationary. 9) Ball A and the blue ball collide simultaneously with ball B, pushing it through the gate. 10) Ball A collides with ball B and ball B goes through the gate. Here, it is unclear whether ball B would have gone through on its own without ball A.

illustrate how they work via the example cases in Figure 3.

Causality Module

The causal knowledge component of our model is based on the CSM (Gerstenberg et al., 2021). The CSM is a quantitative model predicting causal judgments in physical settings (see also Gerstenberg, 2022; Gerstenberg, Peterson, Goodman, Lagnado, & Tenenbaum, 2017; Gerstenberg & Stephan, 2021; Zhou, Smith, Tenenbaum, & Gerstenberg, 2023). The CSM postulates that people make causal judgments by imagining what would have happened in counterfactual situations and comparing those counterfactual outcomes to what happened in the actual situation (Lewis, 1973; Pearl, 2000). The CSM generates counterfactuals by running simulations in an intuitive physics model. Building on prior work, we model intuitive physical thought using a noisy physics simulator (Battaglia,

Did ball A cause ball B to go through the gate?

Figure 4. An example of the billiard ball setting and the objects it contains. The green ball labeled A is the candidate cause, and the grey ball labeled B is the target of the causal interaction. The blue ball and the box are auxiliary objects which can influence the outcome. The brown door can block the exit. If any object contacts the green button below or above, the door will move in the direction of the button that was pressed and come to a stop once it touches the side wall.

Hamrick, & Tenenbaum, 2013; Gerstenberg & Tenenbaum, 2017; Kubricht, Holyoak, & Lu, 2017; Smith et al., in press; Ullman, Spelke, Battaglia, & Tenenbaum, 2017).

Noisy simulators allow us to capture uncertainty about what would have happened in relevant counterfactual situations. When humans simulate how a counterfactual would have played out, various sources of noise affect the accuracy of their simulations (Smith & Vul, 2013). To model this uncertainty, we inject a small amount of Gaussian noise to the direction of the velocity vectors of the objects at the point at which the counterfactual simulation diverges from what actually happened. For example, in Figure 4, if we want to simulate the counterfactual where ball A was not there, we take ball A out of the scenario and run the simulation forward. In the counterfactual, Ball B's path diverges from what happened in the actual situation at the time point where the collision took place. At this point in the counterfactual simulation we begin applying noise to ball B's velocity, reflecting people's uncertainty about what would have happened. The amount of noise in

the physical simulations is determined by a free parameter of the model, θ, the standard deviation of the Gaussian noise distribution.

The CSM posits that people are sensitive to multiple *aspects* of causation. These aspects represent different causal features that have been shown to affect causal judgments. The CSM computes the different aspects of causation by simulating different counterfactual possibilities. Here, we consider three aspects of causation: whether-causation, how-causation, and sufficient-causation.

whether-causation.   Whether-causation corresponds to the notion of counterfactual necessity. This is the traditional counterfactual concept of causation according to which A caused B when both A and B took place and when B would not have taken place without A. The CSM evaluates whether-causation $\mathcal{W}$ by computing the probability that the counterfactual outcome e' in scenario s would have been different from what actually happened e if the candidate cause A had been *removed*.

$$\mathcal{W}(A \rightarrow e) = P(e' \neq e | s, remove(A)). \tag{1}$$

In words, to test if ball A was a whether-cause of ball B going through the gate (e), we stage a counterfactual where we remove ball A from the scene and then simulate what would have happened in that counterfactual simulation. If ball B would have gone through the gate anyway, then we determine that ball A did not make a difference to whether or not the event occurred, so it is not a whether-cause. However, if ball B would not have gone through the gate in the counterfactual, we determine that ball A was indeed a whether-cause of the outcome. This single evaluation yields a binary determination, but as we noted above, people exhibit uncertainty in their counterfactual judgments. To capture this gradation, we compute the probability that ball B would have gone through the gate in ball A's absence by running multiple noisy simulations. The proportion of simulations in which the outcome in the counterfactual situation would have been different from what

actually happened yields a graded measure of whether-causation.

For example, consider scenario 1 in Figure 3. In this scenario, ball B is stationary in the middle of the screen, and ball A comes in from the side and collides with ball B, launching it through the gate. To test for whether-causation in this situation, we remove the candidate cause (ball A) from the scene and then run multiple counterfactual simulations in its absence. In this case, it is very clear that ball B would not have gone through the gate in the counterfactual because it was stationary at the start of the scenario and only picked up momentum after the collision with ball A. After simulating multiple times, the model determines that in scenario 1, ball A has a whether-cause value of 1.0 (the counterfactual is always different from the actual outcome). While the evaluation is very clear in this scenario, it is less clear in scenario 10, for example, where ball B has its own initial momentum. If we ran noisy simulations after removing ball A from this scenario, ball B would go through the gate on some simulations and miss the gate on others. With a noise value of $\theta = 1.0$, the model computes a whether-cause value of $\mathcal{W}(A \rightarrow e) = 0.76$ in this scenario.

how-causation. Counterfactual necessity is an important part of the story of how people make causal judgments. However, it's not the full story. Take for example scenario 2 in Figure 3 where ball A knocks a box that is blocking ball B's path out of the way and then ball B goes through the gate. This is a case of double prevention: ball A prevented the box from preventing ball B from going through the gate. Just as in scenario 1, the presence of Ball A is counterfactually necessary for ball B to go through the gate. However, prior work shows that people often rate candidate causes in double prevention scenarios as less causal than in more standard cases like scenario 1 (Gerstenberg et al., 2021; Henne & O'Neill, 2022; Lombrozo, 2010). Process theorists explain the difference between these two scenarios by appealing to the direct transfer of force from the cause to the target in scenario 1 (e.g. Wolff, Barbey, & Hausknecht, 2010). In scenario 2

there is no direct transfer.[1]

The CSM incorporates information about the causal process by testing for how-causation. Testing for how-causation determines whether a candidate cause made a difference to *how* the outcome came about at a fine level of granularity. Whereas whether-causation assesses difference-making at the level of the outcome event (whether ball B went through the gate or didn't), how-causation is sensitive to the precise details of how that event came about. A candidate cause A, is a how-cause of the fine-grained outcome Δe, if in scenario s, the fine-grained counterfactual outcome $\Delta e'$ would have been different if the candidate cause A had been *changed*:

$$\mathcal{H}(A \to \Delta e) = P(\Delta e' \neq \Delta e | s, change(A)). \tag{2}$$

We define the "fine-grained outcome" as the precise position and time at which ball B passes through the gate, while the *change* operator is implemented as a small perturbation to the initial position of the candidate cause. If this small perturbation leads to a difference in the final position or time at which ball B passes through the gate, the candidate cause is a how-cause.[2] In scenario 1, ball A is indeed a how-cause, the slight change in A results in a slight change in the fine-grained outcome. By contrast, in scenario 2, the perturbation makes no difference. Thus, the notion of how-causation helps us understand why people rate these two scenarios as more or less causal and accounts for that difference with a formal test.

sufficient-causation.   Whether-causation and how-causation express much of the causally relevant information about what happened. However, these two components still fail to capture certain intuitions. A notable set of objections to counterfactual theories of

---

[1]While process theorists can account for the difference between these two cases, they struggle to account for why people feel any inclination to give a causal rating in scenario 2 where there is no direct transfer of force to the target (but see Wolff et al., 2010).

[2]Note that unlike whether-causation, how-causation is binary. We only run a single simulation to determine whether the candidate cause is a how-cause.

causation center around cases of causal preemption (Bunzl, 1980; Hall, 2004; McDermott, 1995; Wolff, 2007). Figure 3 scenario 5 depicts a preemption case. In this scenario, ball A knocks into ball B and sends it through the gate, but even if ball A hadn't been present, the blue ball would have knocked ball B through the gate anyway. Prior work shows that people judge ball A to have caused ball B to go through the gate in this situation even though ball A is not counterfactually necessary for that outcome (Gerstenberg et al., 2021). Whether-causation and how-causation alone cannot explain this pattern.

To capture people's intuition in preemption cases, it seems necessary to include sufficiency (Beckers, 2021; Gerstenberg et al., 2021; Halpern & Pearl, 2005; Icard, Kominsky, & Knobe, 2017; Woodward, 2006). Intuitively, sufficiency tells us whether a candidate cause would have been enough to bring about the outcome "on its own". A counterfactual test for sufficiency simulates whether the candidate cause would have been a whether-cause in the counterfactual contingency in which alternative causes had been removed. For example, if we wanted to know whether ball A was a sufficient-cause for ball B to go through the gate in scenario 5, we would first remove the alternative causes from the scene (in this case the blue ball), and then test in this contingency if ball A would have been a whether-cause. Accordingly, a candidate cause A is a sufficient-cause of outcome e in scenario s if:

$$\mathcal{S}(A \to e) = P(\mathcal{W}(A \to e)|s, remove(\backslash A)). \tag{3}$$

Here, $remove(\backslash A)$ designates the counterfactual operation where we remove all alternative causes from the scene. Once the alternatives are removed, we run the whether-cause test to see whether the candidate cause would have been enough to bring about the outcome on its own. In the situation where there are no alternative causes, the test reduces to a simple test of whether-causation.

We consider an object to be a potential alternative cause if there exists a

counterfactual contingency where it would have been a whether-cause of the outcome. To test this, we simulate each counterfactual contingency, removing every subset of objects (excluding the target ball) and checking whether the candidate alternative was a whether-cause in any of these contingencies.[3] For example, in scenario 5, the blue ball is a whether-cause in the counterfactual contingency where ball A is removed. Thus, we consider the blue ball to be an alternative cause and remove it as part of the sufficient-cause test. If the blue ball hadn't been present, then ball A would have been a whether-cause, thereby satisfying our definition of sufficient-cause.

While this definition of sufficient-causation can help us explain why participants think ball A is a cause in scenario 5, it also makes the counter-intuitive prediction that the blue ball is a sufficient-cause in this situation. If we ran our test on the blue ball, we would remove the alternative cause (ball A) and then find that the blue ball is a whether-cause in this counterfactual contingency. This assessment is problematic: the blue ball clearly didn't cause ball B to go into the gate (Chang, 2009; Gerstenberg et al., 2021). To address this concern, we build on the definition of actual causation in Halpern and Pearl (2005). We further constrain sufficient-causation to check whether the events in the counterfactual contingency match the events that actually happened. The events we consider are the outcome, the obstacle collisions, and objects colliding with either of the green buttons to open or close the gate. We exclude events like balls entering the scene and balls colliding with the walls. Events are defined by the objects involved in them, and not the fine-grained details of their timing. Note that this check for whether events between situations match is asymmetric: events in the counterfactual contingency need to match those in the actual situation, but not vice versa.

To illustrate how this constraint impacts the sufficient-cause test consider scenario 5

---

[3]When checking whether an object is a potential alternative cause, we run the test for whether-causation with a single deterministic simulation. Running multiple simulations across every contingency proved to be prohibitively expensive. In principle, one could run multiple simulations for each contingency to capture uncertainty about which objects qualify as potential alternative causes.

Table 1

Model predictions for scenarios 1–4 shown in Figure 3. a) Aspect values computed for each of the different scenarios. b) Semantic valuations for each of the different causal expressions in those scenarios on the basis of the aspect values. c) Literal listener distributions over scenarios given a particular utterance. These are computed by normalizing the semantic values across scenarios. d) Speaker distributions for a first-level pragmatic speaker. These are computed by renormalizing the distribution of the literal listener across utterances.

a) Aspect Values

Scenario	1	2	3	4
Whether	1.00	1.00	0.00	0.00
How	1.00	0.00	1.00	0.00
Sufficient	1.00	1.00	0.00	0.00

b) Semantic Values

Scenario	1	2	3	4
No Difference	0.00	0.00	0.20	1.00
Affected	1.00	1.00	1.00	0.00
Enabled	1.00	1.00	0.00	0.00
Caused	1.00	0.00	0.00	0.00

c) Literal Listener Distributions

Scenario	1	2	3	4
No Difference	0.00	0.00	0.17	0.83
Affected	0.33	0.33	0.33	0.00
Enabled	0.50	0.50	0.00	0.00
Caused	1.00	0.00	0.00	0.00

d) Speaker Distributions

Scenario	1	2	3	4
No Difference	0.00	0.00	0.33	1.00
Affected	0.18	0.40	0.67	0.00
Enabled	0.27	0.60	0.00	0.00
Caused	0.55	0.00	0.00	0.00

again. Here ball A collides with ball B in the actual situation and also in the counterfactual contingency where we remove the blue ball. Ball A is also a whether-cause in this counterfactual contingency so we determine that it is a sufficient-cause. However, when we run the test for the blue ball, the story is different. In the counterfactual contingency where we remove ball A, the blue ball collides with ball B. This event does not happen in the actual situation. Even though the blue ball is a whether-cause in this contingency, the events that make it so didn't take place in the actual situation, so it doesn't satisfy our definition of sufficiency.

Sample Cases. We illustrate the model computations for each component of the model using the first four cases from Figure 3. Table 1a shows the computed aspect values for these cases. As we've noted above, in scenario 1 ball A is a whether-cause and a how-cause. It is also a sufficient-cause because there are no alternative causes. In scenario 2, ball A is a whether-cause but not a how-cause. It is also a sufficient-cause due

to the fact that again there are no alternatives. In scenario 3, ball A is neither a whether-cause nor a sufficient-cause, but it is a how-cause. And in scenario 4, ball A isn't a whether-cause, a how-cause, or a sufficient-cause.

Semantics Module

We define a semantics that maps from people's causal representation of what happened to causal expressions. We consider four causal expressions: "affected", "enabled", "caused", and "made no difference". In defining the semantics for each of the expressions, we will use $\mathcal{W}$, $\mathcal{H}$, and $\mathcal{S}$ as a shorthand for $\mathcal{W}(A \rightarrow e)$, $\mathcal{H}(A \rightarrow \Delta e)$, and $\mathcal{S}(A \rightarrow e)$, respectively.

"Affected". We define "affected" as

$$\text{AFFECTED}(A \rightarrow e) = \mathcal{W} \vee \mathcal{H} \vee \mathcal{S}. \tag{4}$$

A affected the outcome e if A was a whether-cause, a how-cause, a sufficient-cause, or any combination of the three. "Affected" is the most inclusive causal expression. If the candidate shows any of the different aspects of causation then it affected the outcome. According to this definition, ball A affected ball B's going through the gate in all of the sample scenarios in Figure 3 except scenario 4.

"Enabled". We define "enabled" as

$$\text{ENABLED}(A \rightarrow e) = \mathcal{W} \vee \mathcal{S}. \tag{5}$$

For A to have enabled e it must have either been a whether-cause, a sufficient-cause, or both. For example, in scenario 2 of Figure 3, ball A enabled ball B's going through the gate. It was both a whether-cause and a sufficient-cause (because there were no alternative causes) of the outcome. Scenario 6 demonstrates that including sufficient-causation is important for this causal expression. Here, ball A hits the button that opens the door to the gate shortly before the blue ball would have hit the button. Ball A wasn't a

whether-cause in this case, but there is a strong intuition that ball A enabled ball B's going through the gate.

"Caused".   We define "caused" as

$$\textsc{caused}(A \to e) = \mathcal{H} \wedge (\mathcal{W} \vee \mathcal{S}) \wedge_\sigma \mathcal{M} \wedge_\sigma \mathcal{U}. \tag{6}$$

A caused *e* when it was a how-cause of the outcome, and either a whether-cause or sufficient-cause (or both). In addition to these counterfactual components of the semantics, we further require that the candidate cause *A* was initially moving $\mathcal{M}$ rather than stationary. This requirement is a soft condition as indicated by $\wedge_\sigma$. This means that, even if ball A was stationary at the beginning, it can still be said to have caused the outcome, but the probability of doing so is less than if *A* was moving. The degree of softening is controlled by a parameter σ. Prior work suggests that movement affects people's causal intuitions (Mayrhofer & Waldmann, 2016; White, 2014), and we manipulate the movement of the candidate cause in our stimuli. For example, scenario 7 and scenario 8 are identical in terms of their causal aspects, but in scenario 7, ball A is moving while in scenario 8 it is stationary. As we will see below, participants are more inclined to choose "caused" in scenario 7 than in scenario 8.

Additionally, our definition for "caused" (softly) requires that the candidate cause *A* be unique $\mathcal{U}$. We defined uniqueness in the following way: A contacted B, and no other candidate cause contacted B. This requirement is motivated by the observation that the expression "A caused B" is ambiguous between two senses of "caused". On the one hand, it could be taken to mean that A was *a cause* of the outcome, on the other, it could mean that A was *the cause* of the outcome. Prior work suggests that when multiple causes play a role in an outcome's occurrence, people are inclined to distribute the causal contribution between them (Gerstenberg et al., 2021; Lagnado, Gerstenberg, & Zultan, 2013; Langenhoff, Wiegmann, Halpern, Tenenbaum, & Gerstenberg, 2021; White, 2014; Wu,

Sridhar, & Gerstenberg, 2023). This issue becomes salient in cases like scenario 9. Here

both ball A and the blue ball collide with ball B simultaneously, knocking it into the gate.

Though one could describe this situation by saying "ball A caused ball B to go through the

gate", doing so elides the equally important role that the blue ball played. It seems more

appropriate to say that "A caused B" if A was a unique cause.[4] Like $\mathcal{M}$, $\mathcal{U}$ is a soft

requirement. The degree of softening is also controlled by σ, the same parameter that

determines the level of softening for the movement feature.

"Caused" is the strongest expression in that it has the strictest requirements. A

candidate can only be said to have caused the outcome if it made a difference to how it

came about, and if it was either necessary or sufficient (or both). It is further restricted by

the softer constraints that a cause must be moving and that it be unique.

"Made no difference".   The expression "made no difference" asserts a lack of

causal connection between the candidate cause and the outcome. We define the expression

as the conjunctive negation of each of our causal aspects.

$$\text{NO DIFFERENCE}(A \rightarrow e) = \neg\mathcal{W} \wedge \neg_\nu\mathcal{H} \wedge \neg\mathcal{S}. \tag{7}$$

In words, A made no difference to e if it is not a whether-cause of e, not a how-cause of e,

and not a sufficient-cause of e. The requirement that A not be a how-cause is soft (which

we represent with the soft-not $\neg_\nu$). This softening is intended to capture an ambiguity in

the meaning of "made no difference" that we noted in cases like scenario 3. Here, ball B is

headed toward the gate on its own when ball A comes up behind it and pushes it along.

Ball A is a how-cause because it affected the fine-grained process, but it is neither a

whether-cause nor a sufficient-cause. Even though ball A is a how-cause, it still seems

reasonable here to say that ball A "made no difference" to ball B's going through the gate

----

[4]This same competition does not seem to apply to the other expressions. If we say that "Ball A affected
ball B's going through the gate" it's perfectly fine for the blue ball to have affected ball B as well. The same
seems true for "enabled".

because ball A made no difference to *whether* ball B went through the gate (it only made a difference to *how* it went through). In our experiment, we fit a parameter ν to capture the probability of responding "made no difference" even when the candidate was a how-cause.

Sample Cases.   Table 1b shows how the semantic evaluations of our sample scenarios are shaped by their underlying aspect evaluations. In scenario 1, ball A is a whether-cause, a how-cause, and a sufficient-cause, so it satisfies the definition for "affected", "enabled", and "caused", but not "made no difference". In scenario 2, ball A is a whether-cause and a sufficient-cause. It satisfies the definition for "affected" and "enabled", but because it is not a how-cause, it does not satisfy the definition for "caused". In scenario 3, ball A is a how-cause so it satisfies the definition of "affected" but not "enabled" or "caused". It also weakly satisfies the definition for "made no difference" due to the softening parameter that allows this expression to be applied even when the candidate is a how-cause. The softening parameter in this case was set to 0.2. Finally in scenario 4, ball A does not pass any of the counterfactual tests. As such, the only causal expression it satisfies is "made no difference".

Pragmatics Module

Our semantics links the underlying psychological representation to the meanings of the four different causal expressions. While our semantics tells us when each of these expressions is true, it doesn't tell us how people choose a particular expression in context.[5] This poses a difficulty if multiple expressions are true of the same circumstance. How is a person to choose?

In one of the seminal works of linguistic pragmatics, Grice (1975) illustrates how people talking with one another make sophisticated inferences about what each person intends to communicate that go beyond the literal meaning of the words they use. The underlying principle guiding these inferences is that each person involved in the

---

[5]Technically, the causal expressions themselves do not have truth values. Rather utterances containing the causal expressions have truth values. For the purposes of exposition, we will refer to the truth of causal expressions.

communicative exchange wants to cooperate to help the other person understand. Grice suggests a number of maxims that people generally obey in order to remain cooperative. One of those maxims is a tendency to speak informatively. People generally expect the person they are talking to to tell them as much as they know and is helpful.

A common example of this type of informative communication shows up in the relationship of the words "some" and "all". "Some" and "all" overlap in their meanings. If there are a set of cups on a table that have soda in them, it is true that "all of the cups have soda in them," but it is also true that "some of the cups have soda in them". However, if you ask someone how many of the cups have soda in them and they know the answer, they will tell you that "all of them do" even though saying "some of them do" is strictly speaking true. In order to be cooperative, the speaker not only chooses a true statement, but one that is also as informative as possible. Relatedly, if they only said that "some of the cups have soda in them", you would likely infer that not all of the cups have soda in them even though their sentence is consistent with a world where all the cups have soda in them. This is because you are expecting the speaker to be as informative as they can. You make the added inference that the speaker would have used the more informative utterance if they could have, but since they didn't, it must not be the case. This complex meta-cognitive inference is a classic example of pragmatic reasoning called scalar implicature (Degen, 2015; Hirschberg, 1985).

We can model these types of pragmatic inferences using the rational speech acts framework (RSA, Degen, 2023; Frank & Goodman, 2012; Goodman & Frank, 2016). RSA is a formalization of Gricean communicative theory that posits a speaker and a listener who reason recursively about each other's mental states in order to choose utterances that are both true and informative. Our implementation of the RSA model starts with a literal listener $L0$. Given an utterance with a causal expression, the literal listener infers a distribution over scenarios where any scenario that is consistent with the utterance is

equally probable. Formally,

$$P_{L0}(s|u) \propto \mathcal{M}(s, u). \tag{8}$$

The meaning function $\mathcal{M}$ is our semantics. It takes a scenario s and an utterance u, and returns a semantic value $\in [0, 1]$. P_{L0} is a distribution on scenarios where the semantic values are normalized to sum to one within each utterance. The change in values from Table 1b to Table 1c illustrates how the literal listener transforms the semantic values into a distribution over scenarios. In this example, we limit the set of scenarios over which the listener reasons to scenarios 1–4 from Figure 3. Later, we apply the model to the full set of scenarios in our experiment.

Next, we define a level-1 pragmatic speaker who chooses an utterance so that the literal listener is likely to infer the scenario that the speaker observed. Given a scenario s, the probability that a speaker will choose an utterance u is proportional to the literal listener's assessment of the probability of s given u. Formally,

$$P_{S1}(u|s) \propto P_{L0}(s|u)^{\lambda}. \tag{9}$$

Table 1c and 1d show that this step simply amounts to another round of re-normalization. Whereas the literal listener normalized the semantic values across scenarios (the columns of Table 1), the pragmatic speaker now normalizes the literal listener probabilities across the utterances (the rows of Table 1). Examining the speaker distribution in Table 1d shows an interesting consequence: in scenario 1, the pragmatic speaker favors the expression "caused" over the other two alternatives even though all of them are equally true (see the semantics in Table 1b). This is because "caused" is the most informative expression. It is semantically more restrictive and thus true of a smaller set of scenarios. The hierarchical re-normalization of RSA results in the speaker favoring this more informative utterance, capturing people's tendency to be informative.

The λ parameter in the formula above is a softmax parameter which controls the

optimality of the informative inference. By increasing this parameter, the effect of informativity intensifies. As λ increases, the mass of the distribution concentrates around the most probable response. As λ approaches zero, the distribution becomes increasingly flat, being uniform at $\lambda = 0$. The example in Table 1 assumes $\lambda = 1$.

Recursive reasoning in RSA can be repeated to an arbitrary depth. We can construct a pragmatic listener that reasons about a level-1 pragmatic speaker, and a level-2 pragmatic speaker that reasons about a level-1 listener. Additional levels of recursion increase the effect of informativity, but also impose increased computational costs. In this work, we model participant speakers with a level-2 pragmatic speaker, and participant listeners with a level-1 pragmatic listener.

Sample Cases.   Tables 1c and d illustrate pragmatic inferences for a listener and a speaker. For each utterance, the literal listener (1c) normalizes the semantic values across the set of scenarios. If the utterance is true in multiple scenarios, the literal listener assigns equal probability to each of those scenarios. For example, if the literal listener hears that "Ball A enabled ball B to go through the gate", the literal listener will infer that scenarios 1 and 2 are equally possible because the utterance truthfully describes those scenarios. The literal listener's inference can be modulated by the graded semantic evaluation. For the utterance "Ball A made no difference to ball B's going through the gate", the literal listener assigns most probability to scenario 4, but also applies some probability to scenario 3 because that scenario weakly satisfies the definition for "made no difference" due to the softening parameter.

Table 1d illustrates the inference for the pragmatic speaker. The pragmatic speaker normalizes the probabilities of the literal listener across utterances, resulting in a relative increase in the probability of utterances that are more informative (i.e. true of fewer scenarios). Because the utterance "Ball A caused ball B to go through the gate" is only true in scenario 1, the pragmatic speaker assigns this utterance more probability than the corresponding utterances with "enabled" and "affected". This reflects a pragmatic norm to

choose the most informative utterance that is true in a given scenario.

Experiment 1: Validating Causal Expression Semantics

According to our model semantics, the three causal expressions "caused", "enabled", and "affected" overlap in meaning. However, some of the expressions are more specific than others. They refer to a smaller set of possible situations and therefore are more informative (see Figure 2, semantics module). The most specific and informative expression is "caused", followed by "enabled", and then "affected". In order for "caused" to be true, the conditions for "enabled" and "affected" must also be satisfied and so we say that "caused" implies "enabled" and "affected". Similarly, in order for "enabled" to be true, the conditions for "affected" must also be satisfied, so under our semantics, "enabled" implies "affected".

This scale of implication further grounds our model pragmatics, which predicts that participants will favor more informative expressions when multiple expressions are true of a given situation. Our model predicts that participants will derive scalar implicatures. In a situation in which a listener hears a less specific utterance such as "enabled", they will infer that the more specific utterance, "caused", is not true.

This hypothesis of a scale of specificity and the pragmatic reasoning that it suggests is partially in conflict with the prior literature. Specifically, prior work has proposed that the expressions "caused" and "enabled" are inconsistent. If A caused B it cannot also be the case that A enabled B, and similarly if A enabled B it cannot have caused B. Wolff (2007) makes this claim explicitly, and it also follows from the mental representations that he proposes for "caused" and "enabled". If an agent and a patient relate to each other in the "cause" force configuration, they can't also be in the "enable" force configuration. The two are conceptually as well as semantically inconsistent with one another.

Mental models and causal models yield similar conclusions to a weaker or stronger degree. As we discuss above, the mental model for "caused" can be paraphrased as saying that A is sufficient but not necessary for B, while the mental model for "enabled" can be

paraphrased as saying A is necessary but not sufficient for B.[6] If "caused" implies sufficiency and "enabled" precludes it, and similarly if "enabled" implies necessity while "caused" precludes it, then once again the two concepts are inconsistent.

Causal models, too, suggest that "caused" and "enabled" are inconsistent. We can again see this by analyzing what these definitions imply about necessity and sufficiency. The causal model for "caused" (see Figure 1B) implies that A is necessary *and* sufficient for B. A is the only causal factor determining B, so if A is active B will be as well, and if A is not active B will not be active either. On the other hand the causal model for "enabled" implies that A is necessary but not sufficient for B. A and the auxiliary variable X determine B in a conjunctive relationship. Both must be active for B to be active, so A is necessary. But A's activity alone is not enough for B to be active, so A is not sufficient. Because "caused" implies sufficiency but "enabled" precludes sufficiency, again these concepts are inconsistent.

These prior accounts have less to say about the relationship between "affected" and the other causal expressions. None of the prior models provide explicit mental representations for this expression, however Wolff, Klettke, Ventura, and Song (2005) comment about its semantic relationship to "caused" and "enabled". They suggest "affected" statements exist on a hierarchy of causal expressions. They are the lowest level of that hierarchy and thereby very inclusive in their meaning and compatible with many stronger forms of causal expressions, including expressions containing periphrastic causatives like "caused" and "enabled". This view seems largely compatible with the claims we have made so far about the relationship between "affected", "caused", and "enabled".

In summary, though our account of the meaning of "affected" seems to be aligned with Wolff et al.'s (2005) statements on its meaning, our claims about the relationship

---

[6]Mental Models Theory actually provides a mental model for the verb "allow" rather than "enable". Though subtly different, these two verbs are often used interchangeably in the literature, and we consider them to have relatively aligned meaning for the purposes of making comparisons between models in this paper. See Goldvarg and Johnson-Laird (2001) for some additional discussion of the distinction between the two.

between "caused" and "enabled" conflict with the semantics proposed by all previous accounts. In Experiment 1 we validate the general hierarchy of specificity we proposed in our semantics, and show that participants' intuitions about relationships among these causal expressions align with our proposal in contrast to the claims of these prior models. Experiment 1 includes three studies: an initial study where we tested the baseline acceptability of our causal expressions in a variety of sentence frames, a second study to assess whether participants' intuitions about the causal expressions are consistent with our semantics, and a third study to assess whether participants cancel implicatures as predicted by our pragmatic account. For the second and third study, we offer comparisons of our own model's predictions against the predictions of prior accounts.

Experiment 1A: Norming study

If saying that "The new technology caused the change." implies that it enabled and affected it, it must at a minimum be acceptable to use all three of these verbs in the same sentence frame. To test our hypotheses, we need a collection of sentence frames where all three causal expressions are acceptable. To collect these stimuli, we first ran a norming study that examined the acceptability of using the different causal expressions in a set of sentence frames. The sentence frames that participants rated acceptable across all the expressions in this norming study then served as the base for the stimuli in our follow-up experiments where we tested our primary hypotheses.

Methods

All experiments were approved by Stanford's IRB (#48665). Experiment 1A was developed and deployed using the jsPsych experiment library (De Leeuw, 2015). We pre-registered our data-collection paradigm and analysis plans on the Open Science Framework: `https://osf.io/kx5fg`. The data, study materials, and analysis code for all experiments in this paper are available here:

`https://github.com/cicl-stanford/causal_language`

Table 2

Sample sentence frames from the norming study. The top ten sentence frames had median ratings above the midpoint of the scale for all three causal expressions. These sentence frames were used to construct stimuli in the follow-up experiments. The bottom three are a sample of the sentence frames that were excluded. For these sentence frames, the median response for at least one expression was at the midpoint of the scale or below.

Included Sentence Frames
1. The dry weather _____ the wild fire.
2. The CEO's decision _____ the outcome.
3. The new technology _____ the change.
4. The Sacklers' greed _____ the opioid epidemic.
5. The sunny weather _____ the tree's growth.
6. More stipends _____ the increase in student admissions.
7. The sun _____ the drying of the clothes.
8. Metastasis _____ cell growth.
9. Diversification _____ new monetary policies.
10. The algae buildup in the ocean _____ the migration of certain species of fish.

Examples of Excluded Sentence Frames
1. The collapse of Lehman Brothers _____ the financial crisis.
2. Janelle's working hard _____ her success.
3. Turning off the life support _____ the patient's death.

Participants.   We recruited 51 participants (*age*: M = 35, SD = 12, *gender*: 25 female, 23 male, 3 non-binary, *race*: 37 White, 6 Asian, 6 Black/African-American, 2 other) via Prolific. We restricted selection to participants who had completed at least 10 previous experiments and have an overall 95% experiment approval rating. All participants were fluent in English and based in the United States (we use the same inclusion criteria for the subsequent studies in Experiment 1). We excluded one participant who failed to pass the attention check, leaving a total of 50 participants for analysis. Participants were paid $2 for 10 minutes of work.
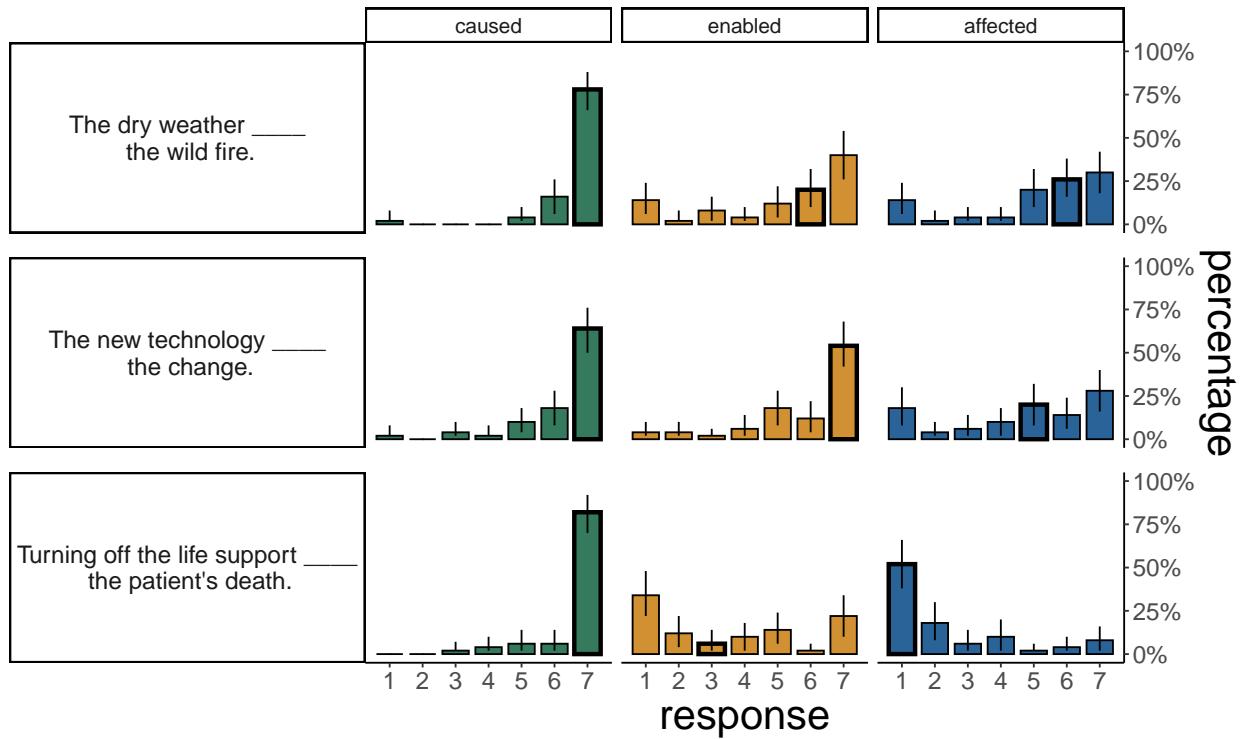
Stimuli.   We developed a set of 20 sentence frames to test for acceptability with each of our causal expressions. For each sentence frame, we created three items by substituting in each of the three causal expressions. In total, we had 60 stimuli sentences. We aimed to collect a set of sentences that covered a wide variety of causal systems, and

the full set of stimuli included scenarios across domains such as medicine, weather, and finance. Table 2 shows some examples, and Appendix A shows the full set of sentence frames. We also included six attention checks, two for each causal expression, which were designed to be either obviously acceptable or obviously unacceptable. Participants who failed to answer on the correct side of the scale for more than two attention checks were excluded from analysis.

Procedure.   We instructed participants that they would see a series of 66 sentences and rate their acceptability. We provided minimal guidance to the meaning of acceptability encouraging participants to trust their own intuitions of what sounds "natural". We provided participants with an example that we labeled as acceptable ("Working long hours caused Pat to feel tired.") and an example that we labeled unacceptable ("Working long hours affected Pat to feel tired."). After reading these instructions, participants proceeded to the main task. Items were presented one by one in randomized order and led by the prompt "Is this an acceptable English sentence?". Participants provided ratings on a 7-point Likert scale with the endpoints labeled "definitely no" and "definitely yes", and the midpoint labeled "unsure". Participants had to provide a judgment on each item to continue.
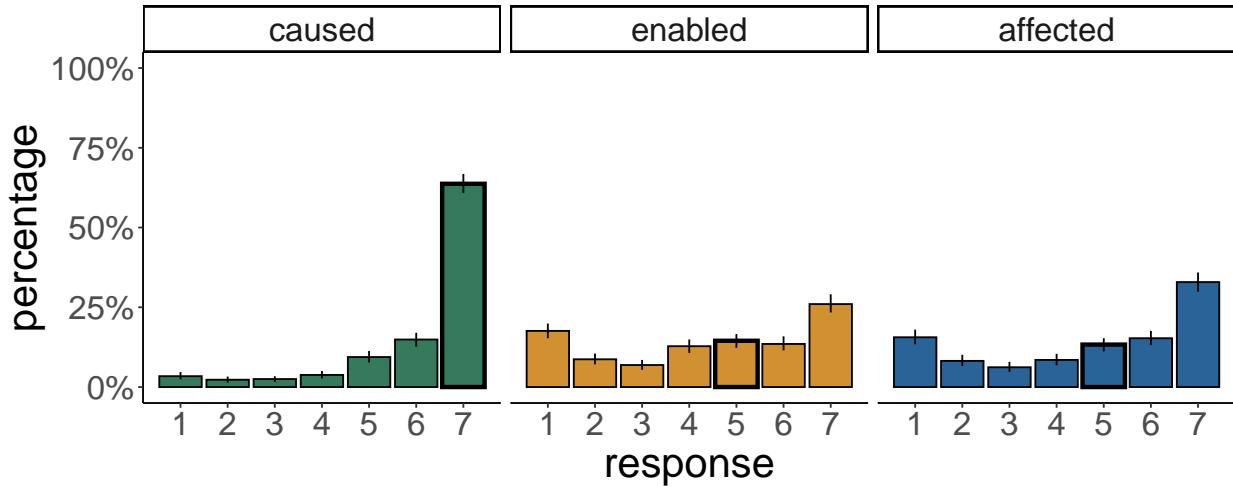
Results and Discussion

Our aim was to collect a set of 10 sentence frames for which all three causal expressions were acceptable. We defined acceptability as a median rating of 4 (the scale midpoint) or higher. 15 of 20 frames met this threshold. In our exploratory analysis, we found that 10 sentence frames had a median acceptability of 5 or higher for all three causal expressions. We selected these 10 sentence frames as our stimuli for our follow-up experiments. Figure 5 shows histograms of acceptability ratings for three example sentence frames. In the top two frames, participants found all three expressions acceptable. In the bottom sentence frame, only "caused" was rated acceptable. Therefore, we didn't include this sentence frame in subsequent experiments.

Figure 5. **Experiment 1A**. Sample sentence frames considered for inclusion in our experiment with histograms for each causal expression showing participant ratings of acceptability. Median responses are indicated by a bold outline. The top two sentences were included because participant median ratings on all three verbs were above the midpoint of the scale. The bottom sentence was excluded because participant ratings for both "enabled" and "affected" fell below the threshold. Error bars represent 95% bootstrapped confidence intervals.

In general participants found items with "caused" more acceptable than items with "enabled" or "affected" for our sentence frames. Figure 6 shows histograms of participant responses for each verb, aggregated across items. The histogram for "caused" skews strongly toward the acceptable end of the scale. On the other hand the histograms for "enabled" and "affected" show more variance, though still tend overall toward acceptability. Considering individual frames, the median rating for "caused" was above the midpoint of the scale on all frames. The rating for "enabled" was above the midpoint in 12 out of 20 frames. The median rating for "affected" was above the midpoint in 14 out of 20 frames.

Figure 6. **Experiment 1A**. Overall responses for each verb aggregated across items. In general, sentence frames with "caused" were rated highly acceptable. Sentence frames with "enabled" and "affected" showed more variance in the distribution of responses. *Note*: Median responses are indicated by a bold outline. Error bars represent 95% bootstrapped confidence intervals.

Experiment 1B: Semantic Relations of Causal Expressions

The stimuli collected in Experiment 1A provide us with a set of sentences to test the semantics of our causal expressions. We want to know whether people's intuitions about the relationships among the expressions align with the structure of the semantics module as shown in Figure 2. When people say that "A caused B", does that imply that "A enabled B" and "A affected B"? Similarly when people say that "A enabled B", does that imply that "A affected B"? While these implications from stronger to weaker should hold if our semantics is true, the reverse is not the case. It should be acceptable for A to enable or affect B without causing it. The implication is uni-directional.

To test this prediction, we augmented our stimuli from the previous experiment with an additional "but it didn't _____ it" clause. For example, the sentence frame "The sunny weather _____ the tree's growth." became "The sunny weather _____ the tree's growth, but it didn't _____ it."[7] To create our stimuli, we substituted each pair of expressions into the

_____

[7]Shibatani (1976) uses a similar construction to illustrate the more restrictive meaning of lexical causatives relative to the periphrastic causative "cause".

blanks in these sentence frames. When the order of the expressions goes from a more specific verb to a less specific verb (e.g. "The sunny weather caused the tree's growth, but it didn't affect it."), we hypothesized that participants would find the sentence unacceptable. This is because, under our semantics, sentences like these express a contradiction. If "caused" implies "affected", then saying that "A caused B, but it didn't affect B" implies that "A affected B, but it didn't affect B". On the other hand, we hypothesized that the reverse ordering of the expressions, where the less specific expression comes first followed by the more specific expression, should be acceptable (e.g. "The sunny weather enabled the tree's growth, but it didn't cause it."). According to our semantics there are possible scenarios where the more specific expression (e.g. "caused") is not true while the less specific expression is (e.g. "enabled").

Methods

Experiment 1B was developed and deployed using the jsPsych experiment library (De Leeuw, 2015). We pre-registered our data-collection paradigm and analysis plans on the Open Science Framework: `https://osf.io/2un9v`.

Participants.   We recruited 55 participants (*age*: $M = 42, SD = 14$, *gender*: 31 female, 24 male, *race*: 46 White, 4 Asian, 3 Black/African American, 2 other) online using the Prolific platform. We excluded 2 participants that failed to pass the attention check. Participants were paid $2 for 10 minutes of work.

Stimuli.   We took all ten of the sentence frames collected in our preliminary study and augmented them with the "but it didn't _____ it" clause as described above. To create our experiment items, we permuted each pair of causal expressions in each frame, leading to a total of 60 items. Table 3 shows six sample items for this experiment on the top of the table. In addition to our primary experimental items, we included the same six attention checks as we had in the preliminary experiment. As in the previous experiment, participants who failed more than two of these attention checks were excluded from analysis.

Table 3

Experiment 1. *Example items from the Semantics Experiment (Experiment 1B) and the Pragmatics Experiment (Experiment 1C). In the Semantics Experiment, items with verb orderings that are contradictory under our semantics are presented on the left, and corresponding items that are consistent with our semantics are shown on the right. In the Pragmatics Experiment, items that use the "in fact" phrasing to cancel the implication are shown on the left, and corresponding items that use the "in fact" phrasing to cancel the implicature are shown on the right.*

Semantics Experiment (Experiment 1B)	
Contradictory Ordering	*Non-Contradictory Ordering*
The Sackler's greed caused the opioid epidemic, but it didn't enable it.	The Sackler's greed enabled the opioid epidemic, but it didn't cause it.
The sunny weather caused the tree's growth, but it didn't affect it.	The sunny weather affected the tree's growth, but it didn't cause it.
Metastasis enabled cell growth, but it didn't affect it.	Metastasis affected cell growth, but it didn't enable it.
Pragmatics Experiment (Experiment 1C)	
Implication Cancellation	*Implicature Cancellation*
The dry weather caused the wild fire, in fact it enabled it.	The dry weather enabled the wild fire, in fact it caused it.
The CEO's decision caused the outcome, in fact it affected it.	The CEO's decision affected the outcome, in fact it caused it.
The new technology enabled the change, in fact it affected it.	The new technology affected the change, in fact it enabled it.

Procedure. The procedure for this experiment was similar to that of Experiment 1A. We instructed participants that their task was to rate the acceptability of 66 sentences. Again, we informed participants that we were interested in their intuitions of whether or not each sentence seemed "natural". We provided two example sentences with the same form as our experimental items. In one sentence a more specific verb followed a less specific verb, and in the other sentence the reverse was true. Unlike in the preliminary experiment, we did not indicate whether either of these example sentences was acceptable or not. Participants proceeded to provide ratings on the same Likert scale as in Experiment 1A. Again, the item order was randomized.

Table 4

Summary of the acceptability predictions for our model (CSM) and the models from the prior literature. ✔ *indicates that the model predicts the ordering will be acceptable.* ✗ *indicates the model predicts the ordering will be unacceptable. Mental Model Theory and Causal Model Theory do not make any predictions about "affected" and so acceptability predictions cannot be provided. In general, where the models do make predictions those predictions align. The key difference is in the "caused" to "enable" ordering where our overlapping semantics predict non-acceptability and the inconsistent semantics of the prior models predict acceptability.*

Ordering	CSM	Mental Model	Causal Model	Force Dynamics
affected → enable	✔	NA	NA	✔
affected → cause	✔	NA	NA	✔
enabled → affect	✗	NA	NA	✗
enabled → cause	✔	✔	✔	✔
caused → affect	✗	NA	NA	✗
caused → enable	✗	✔	✔	✔

Hypotheses. We pre-registered the confirmatory hypothesis that items where the less specific verb preceded the more specific verb (affected → caused, affected → enabled, enabled → caused) would be more acceptable overall than the corresponding items with the reversed orders (caused → affected, enabled → affected, caused → enabled). We broke down this hypothesis into three sub-hypotheses, one for each pair of causal expressions. For example, we predicted that enabled → caused would be more acceptable than caused → enabled.

As an exploratory analysis, we also sought to examine whether our data more closely aligned with the overlapping semantics of "caused" and "enabled" that we propose in our model, or the inconsistent semantics suggested in prior work. What predictions do the "overlapping semantics" account versus the "inconsistent semantics" account make? For sentences like "The dry weather enabled the wild fire, but it didn't cause it", both semantics predict that participants will rate these sentence frames with higher acceptability, though for different reasons. In the inconsistent semantics, these types of sentences are acceptable because the two concepts can't truthfully describe the same
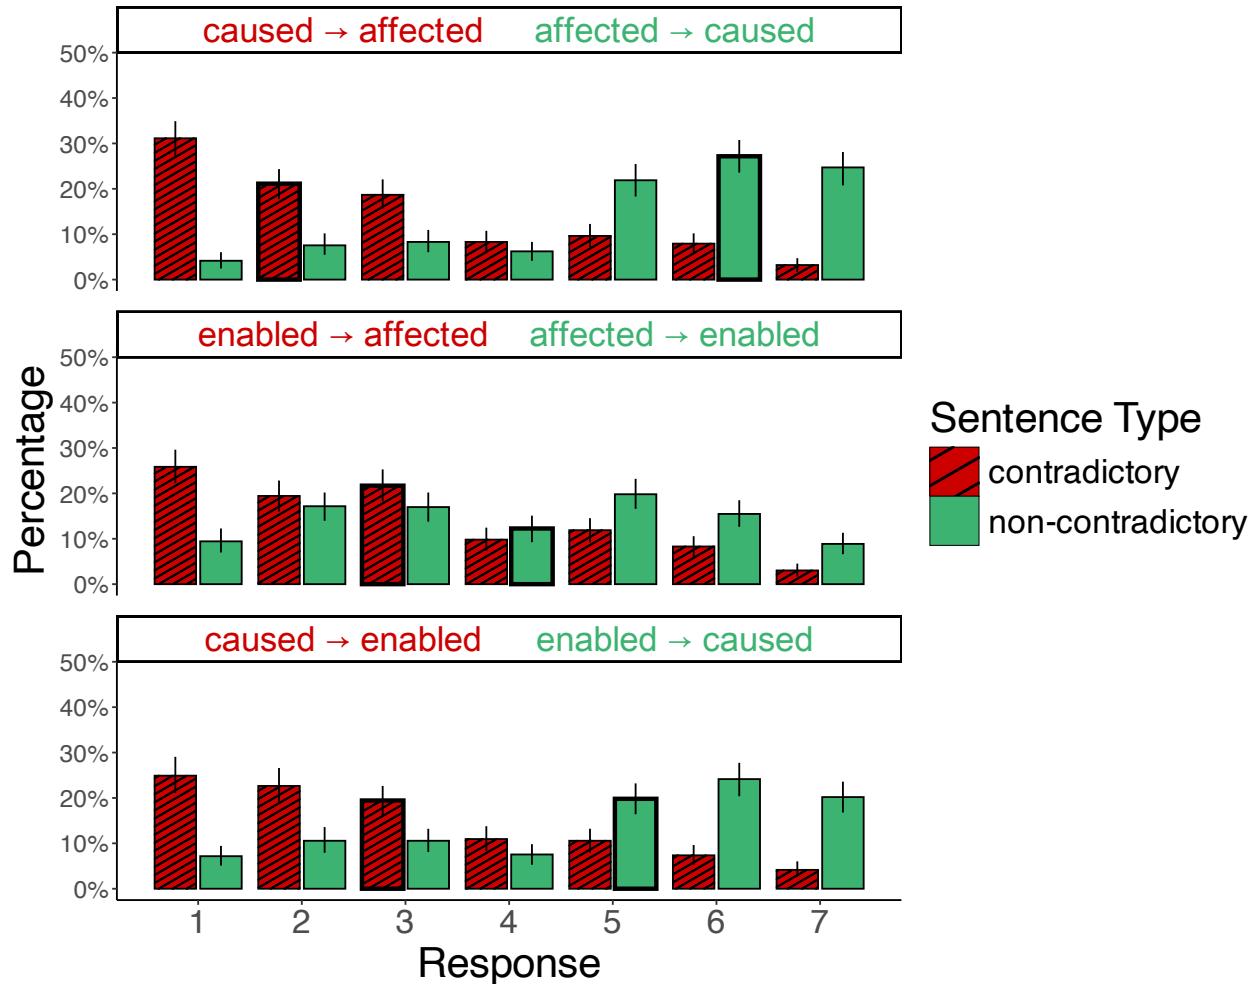
phenomenon. If something "enabled", it did not "cause" and vice versa. In the semantic overlap approach, these types of sentences are acceptable because it is possible to satisfy the semantics for the less specific utterance, "enabled", without satisfying the semantics for more specific utterance, "caused". This semantic position corresponds to the middle rung of the Venn diagram in Figure 2.

The two approaches come apart in their predictions for sentence frames with the reverse verb ordering. For sentences like "The dry weather caused the wild fire, but it didn't enable it", the inconsistent semantics predicts that participants will rate this sentence frame as acceptable. Under this theory, because the two concepts are non-overlapping, it is once again appropriate to endorse a statement where one causal relation obtains while the other does not. But in the semantic overlap approach, this is not the case. By satisfying the semantics for "caused", you have also satisfied the semantics for "enabled" (see the center rung of the Venn diagram in Figure 2). This approach predicts that participants will in general rate sentences with this verb ordering as unacceptable because they are contradictory.

Table 4 illustrates the qualitative acceptability predictions for the different models on each verb ordering. Mental Model Theory and Causal Model Theory only provide predictions for orderings with "caused" and "enabled". As we note, Force Dynamics Theory also suggests a semantics for "affected" which is compatible with "enabled" and "caused". This semantics aligns with our own semantics for "affected", and so the predictions for the orderings involving "affected" are the same between these two models.

Results

Confirmatory analyses.   Figure 7 shows participants' responses. Each row shows the results for a pair of expressions. The red striped bars represent responses on orderings that were contradictory under our semantics, and the green bars represent responses on orderings that were consistent with our semantics. In general, the green bars skew to the right of the red bars, indicating that participants found items which are consistent with our

Figure 7. **Experiment 1B**. Histogram of the overall acceptability of contradictory items and non-contradictory items in the semantics experiment. Median ratings are indicated by bold bar outline. Across all pairings, contradictory items are overall less acceptable than non-contradictory items, and the median rating for the contradictory items is below the midpoint of the scale for all pairings of causal expressions. For the non-contradictory items, the median rating of "affected" → "caused" and "enabled" → "caused" is above the midpoint of the scale, while for "affected" → "enabled" the median is at the midpoint of the scale. Error bars represent bootstrapped 95% confidence intervals.

semantics overall more acceptable than items which are contradictory under our semantics. For example, participants rated items like "The dry weather caused the wild fire, but it didn't enable it." less acceptable than the reversal "The dry weather enabled the wild fire, but it didn't cause it.".

To test our confirmatory hypotheses, we conducted a Bayesian regression analysis.

Table 5

Summary of the confirmatory hypothesis test results for Experiment 1B. For each pairing of verbs, we computed the contrast distribution by subtracting the posterior distribution of the verb ordering that went from the stronger verb to the weaker verb from the corresponding distribution for the ordering from weaker verb to the stronger verb (e.g. subtracting the posterior for the caused → affected ordering from the posterior for the affected → caused ordering). The 'Posterior Estimate' column represents the mean of the resulting contrast distribution. The middle column reports the lower bound of the credible interval. 95% of the contrast distribution lies above this bound. The 'Posterior Probability' column reports the proportion of samples from the posterior density favoring the hypothesis that the contrast distribution lies above zero. For all pairings, all samples spoke in favor of the hypothesis indicating very strong evidential support.

Verb Pairing	Posterior Estimate	Credible Interval Lower Bound	Posterior Probability
AFFECTED-CAUSED	1.47	1.34	~ 1
ENABLED-CAUSED	1.09	0.97	~ 1
AFFECTED-ENABLED	0.62	0.50	~ 1

We fit a hierarchical Bayesian ordinal regression predicting participant responses from the verb ordering in the given item. Verb ordering was represented as a six-level factor, one level for each permutation of our causal expressions (enabled → affected, caused → enabled, etc.). Additionally, we included random intercepts for participants and sentence frames. All Bayesian regressions in this and subsequent experiments were fit using the `brms` package (Bürkner, 2017) in the programming language `R` (R Core Team, 2019) assuming default priors for all parameters.

We computed linear contrasts on the levels of the verb ordering factor of the fitted regression model to test our hypotheses. For each causal expression pairing, we subtracted the posterior estimate of the parameter on the ordering where the first verb implied the second from the posterior estimate for the reverse ordering. For example, we took the distribution representing the parameter estimate for the verb ordering caused → affected and subtracted it from the distribution representing the parameter estimate for the reverse ordering affected → caused. We considered each sub-hypothesis to be confirmed if 95% of the resulting contrast distribution was above zero.

Table 5 shows the results of our confirmatory hypothesis tests. All hypotheses were

confirmed. 95% of the posterior density for each contrast distribution lies above zero, indicating their is a credible difference between the orderings for each verb pairing.

Exploratory analyses. What does our data say about whether the semantics for "caused" and "enabled" are inconsistent or overlapping? The key comparison is illustrated in the "caused"/"enabled" row of Figure 7. As both accounts predict, participants rated sentences that go from "enabled" to "caused" as overall acceptable. The median rating for these sentences indicated by the bold outline on the bar was at 5, above the midpoint of the scale. Critically, for the reverse ordering, sentences that go from "caused" to "enabled", participants seemed to find these sentences overall unacceptable. The median rating for these sentences was at 3, below the midpoint of the scale. Contrary to the predictions of the inconsistent semantics, and aligned with the predictions of our overlapping semantics, the data suggests that participants judge sentences unacceptable that start with the more specific expression but then deny the less specific one.

Discussion

The results of our analysis generally confirm our model semantics. Overall, participants rated sentences that implied a contradiction under our semantics as unacceptable and sentences that were consistent as acceptable. This suggests that, in line with the scale of specificity that we hypothesized, participants believe that "caused" implies "enabled", and that "enabled" implies "affected".

One notable exception is the finding that the median rating for sentences where "affected" preceded "enabled" was at the midpoint of acceptability. There was substantial variance in participants' acceptability ratings for this ordering (see Figure B1). For some items, such as "The CEO's decision affected the outcome, but it didn't enable it.", the distribution of responses skews toward acceptability. For other items, such as "The sunny weather affected the tree's growth, but it didn't enable it.", the distribution skews toward non-acceptability. But even though participants were generally more uncertain about sentences that went from "affected" to "enabled", they were confident that sentences with

the reverse ordering were unacceptable (e.g. "The dry weather enabled the wild fire, but it didn't affect it.").

Participants' judgments for sentences containing the expressions "caused" and "enabled" seem to suggest a semantics where the meanings of these two expressions overlap. Participants generally rated sentences like "The dry weather caused the wildfire, but it didn't enable it" as unacceptable. This is consistent with our model semantics that claims that "caused" implies "enabled" and so sentences of this form are contradictory. This finding contrasts with an account that suggests the meanings of "caused" and "enabled" are inconsistent. Under such an account, a sentence like the one above is perfectly acceptable, causing doesn't imply anything about enabling.

Experiment 1C: Pragmatics of Causal Expressions

Next, we examine whether people exhibit the pragmatic behaviors we would expect based on our model semantics. As we noted above, when semantic scales have this type of informative hierarchy, they give rise to scalar implicatures. If a speaker uses a weaker verb on the scale, a listener will generally infer that the speaker doesn't think the stronger verb is true of the situation being described.

A standard method to test for whether a particular expression is an implicature is to see whether it can be *cancelled* (Grice, 1975; Mayol & Castroviejo, 2013). Recalling the example of "all" and "some", if a speaker tells a listener "Some of the cups have soda in them.", the listener will likely infer that "Not all of the cups have soda in them.". However, the speaker can cancel this implicature in their statement by adding an additional clause: "Some of the cups have soda in them, in fact all of them do." The "in fact" clause adds additional information that is consistent with, but stronger than what was said before (Matsumoto, 1997). In contrast, if one tries to cancel the implication from "all" to "some" the result is less natural: "All of the cups have soda in them, in fact some of them do." Here, the "in fact" clause offers redundant information. Because stating that "all the cups have soda in them" implies that "some of the cups have soda in them", the "in fact" clause

is merely repeating something that was already communicated.

The bottom half of Table 3 illustrates the contrast between sentences that use the "in fact" locution to cancel an implication or an implicature. Under our semantics, the sentences that cancel implicatures should be judged acceptable, while the sentences that cancel implications should appear unnatural and redundant.

Methods

Experiment 1C was developed and deployed using the jsPsych experiment library (De Leeuw, 2015). We pre-registered our data-collection paradigm and analysis plans on the Open Science Framework: `https://osf.io/ak5yd`.

Participants. We recruited 54 participants (*age*: M = 34, SD = 13, *gender*: 27 female, 25 male, 2 Non-binary *race*: 37 White, 6 Asian, 5 Black/African American, 1 American Indian/Alaska Native, 5 other) on the Prolific platform. We excluded 3 participants who failed to pass an attention check, leaving us with a total of 51 participants in our analysis. Participants were paid $2 for 10 minutes of work.

Stimuli. As in the previous experiment, we took the set of ten sentence frames that we collected in Experiment 1A and augmented them. This time, instead of adding a "but it didn't _____ it" clause, we added an "in fact it _____ it" clause. For example, the preliminary frame "The CEO's decision _____ the outcome." became "The CEO's decision _____ the outcome, in fact it _____ it." To create our full set of items we again permuted each pair of causal expressions in each frame, leading to a total of 60 items. Table 3 shows six sample items at the bottom of the table. As in the previous two experiments, we included the same attention check items, with the same criteria for exclusion.

Procedure. The procedure for this experiment was nearly identical to that of the previous study. Participants saw the same instructions except the two sample sentences were replaced by a pair of examples using the "in fact" locution. The only other difference was the items themselves, which were constructed as explained above.

Table 6

Summary of the acceptability predictions for our model (CSM) and the models from the prior literature on "in fact" sentences with the given verb ordering. ✔ *indicates that the model predicts the ordering will be acceptable.* ✗ *indicates the model predicts the ordering will be unacceptable. No prediction is indicated by NA. Again we see the models mostly align in their predictions when they make predictions. Here the key difference is in the "enabled" to "caused" ordering where our model's overlap semantics predicts acceptability and the other models' inconsistent semantics predict unacceptability.*

Ordering	CSM	Mental model	Causal model	Force dynamics
affected → enable	✔	NA	NA	✔
affected → cause	✔	NA	NA	✔
enabled → affect	✗	NA	NA	✗
enabled → cause	✔	✗	✗	✗
caused → affect	✗	NA	NA	✗
caused → enable	✗	✗	✗	✗

Hypotheses

We pre-registered the hypothesis that sentences where the "in fact" clause cancelled an implicature (affected → caused, affected → enabled, enabled → caused) would be overall more acceptable than sentences where the same clause attempted to cancel an implication (caused → affected, enabled → affected, caused → enabled). For example, sentences like the ones on the right side of the bottom half of Table 3 would be more acceptable than sentences like the ones on the left. Once again we broke this down into three sub-hypotheses, one for each utterance pair.

As in Study 2, we were also interested whether the data support a semantics for "caused" and "enabled" that is inconsistent or overlapping. What do the different semantics predict for "in fact" sentences with these two causal verbs? For sentences like "The new technology caused the change, in fact it enabled it", both semantics predict that participants will rate the sentence with low acceptability. For the inconsistent semantics, this type of sentence is unacceptable because the two concepts don't overlap. The "in fact" clause adds information that is consistent with and stronger than the initial clause (Matsumoto, 1997), but if the concepts are inconsistent to begin with, then connecting
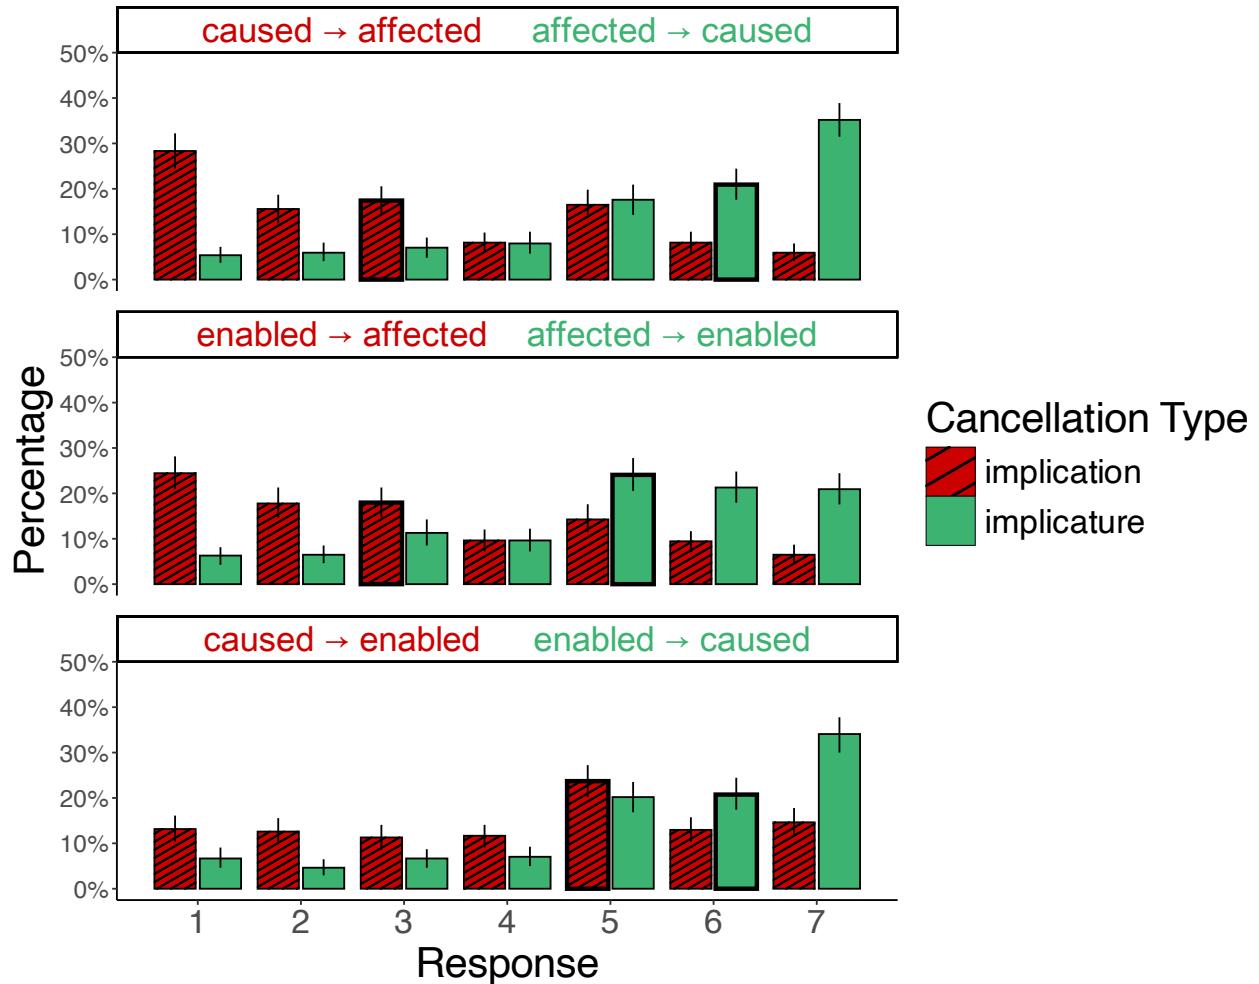
them with this construction should violate people's intuitions about the meaning of the verbs. Similarly for the overlapping semantics, we predict that participants will rate these sentence frames with low acceptability. In this case however, this is because of the position in the semantic hierarchy. The verb ordering that goes from "caused" to "enabled" starts with the more informative verb and goes to the less informative one. It repeats information that is already implied by the first verb. Because of this redundancy, we expect participants to rate these sentences as less acceptable.

The two semantics come apart in their predictions of the reverse ordering. For sentences like "The new technology enabled the change, in fact it caused it", an inconsistent semantics predicts that participants will rate this as unacceptable for the same reasons as before. Because the two concepts are inconsistent with one another the "in fact" construction is inappropriate, it assumes the meanings of the verbs are consistent. On the other hand, the overlap semantics predicts that participants will rate this kind of sentence acceptably. Because "caused" is consistent with and stronger than "enabled", the "in fact" construction is appropriate in this case. This is the logic underlying the cancellation test for the scalar implicature.

Table 6 illustrates the qualitative acceptability predictions for the different approaches with the different verb orderings. Once again, Mental Model Theory and Causal Model Theory don't provide predictions for sentences with "affected", while the predictions of Force Dynamics Theory align with our model predictions. The critical contrast is the "enabled" to "caused" ordering.

Results

Confirmatory analyses.   Figure 8 shows participants' responses. Again, each row represents a particular verb pair. Here, the red striped bars represent items that cancelled an implication, while the green bars represent items that cancelled an implicature. Again, we see that the green bars consistently skew to the right of the red bars. Overall, participants rated items that cancelled implicatures more acceptable than items that

Figure 8. **Experiment 1C**. Histogram of overall acceptability of implicature cancellations vs. implication cancellations for each causal expression pairing. Median ratings are indicated by bold bar outline. Across all pairings, the implicature cancellations are more acceptable than implication cancellations, and the median rating for the implicature cancellations is always above the midpoint of the scale. For the implication cancellations, the median rating of "enabled" → "affected" and "caused" → "affected" is below the midpoint of the scale, while for "caused" → "enabled" the median is above the midpoint of the scale. Error bars represent bootstrapped 95% confidence intervals.

cancelled implications. For example, participants were more likely to give high ratings to items like "The sunny weather affected the tree's growth, in fact it caused it." than items like "The sunny weather caused the tree's growth, in fact it affected it.".

As in the previous experiment, we fit a hierarchical Bayesian ordinal regression predicting participant responses from verb ordering, and then computed linear contrasts on

Table 7

Summary of the confirmatory hypothesis test results for Experiment 1C. For each pairing of verbs, we computed the contrast distribution by subtracting the posterior estimate of the verb ordering that went from the stronger verb to the weaker verb from the corresponding estimate for the ordering from weaker verb to the stronger verb (e.g. subtracting the posterior for the caused → affected ordering from the posterior for the affected → caused ordering). The 'Posterior Estimate' column represents the mean of the resulting contrast distribution. The middle column reports the lower bound of the credible interval. 95% of the contrast distribution lies above this bound. The 'Posterior Probability' column reports the proportion of samples from the posterior density favoring the hypothesis that the contrast distribution lies above zero. For all pairings, all samples spoke in favor of the hypothesis indicating very strong evidential support.

Verb Pairing	Posterior Estimate	95% of posterior density above	Posterior Probability
AFFECTED–CAUSED	1.37	1.21	~ 1
AFFECTED–ENABLED	0.96	0.83	~ 1
ENABLED–CAUSED	0.73	0.57	~ 1

the levels of the verb ordering to test our hypotheses. We subtracted posterior estimates for the three verb orderings where the stronger verb preceded the weaker one from the posterior estimates of the corresponding verb orderings where the weaker verb preceded the stronger one. We tested each pair of causal expressions individually and considered the hypothesis to be confirmed if 95% of the contrast distribution was above zero. Our ordinal regression included random intercepts for participants and sentence frames.

We also included a control in the regression for the raw acceptability of the different causal expressions in each frame. As we noted in the discussion of Experiment 1A, the three causal expressions varied in their acceptability with "caused" being the most acceptable. We hypothesized that independent from any consideration of the implicature, the switch in acceptability from the less acceptable response to the more acceptable one could drive participant responses and confound our effect. For example, in the item pair, "The new technology affected/caused the change, in fact it caused/affected it", participants might overall rate the affected → caused ordering more acceptable than the caused → affected order merely because "The new technology caused the change." is more acceptable than "The new technology affected the change.", and the "in fact" locution

corrects from the less acceptable one to the more acceptable one.

To control for this possibility we computed an acceptability difference predictor for each item. We took the median acceptability scores for both verbs in that item from the norming study, and subtracted the acceptability of the first verb from the second. We computed this score for each item and included it as another predictor alongside the verb-ordering itself.[8]

Table 7 displays the results of our confirmatory hypothesis tests. Again, we see that 95% of each contrast distribution lies above zero, indicating there is a credible difference between each verb order in line with our pre-registered hypotheses.

Exploratory analyses. How do these data bear on the question of whether the semantics for "caused" and "enabled" are overlapping or inconsistent? The key comparison is illustrated on the "caused"/"enabled" row of Figure 8. Contrary to the predictions of both accounts, participants rated sentences that go from "caused" to "enabled" as overall acceptable. The median rating was 5, above the midpoint of the scale, though notably this particular response distribution showcases broad variation. But because the models are aligned in their prediction for this ordering, this failure doesn't help us adjudicate between them.

The critical comparison is the "enabled" to "caused" ordering. Here, participants rated sentences with this verb ordering overall acceptable. The median rating was 6, well above the midpoint of the scale. This pattern aligns with the predictions of the semantic overlap account, which suggests this verb ordering is an appropriate implicature cancellation, but contrasts with the predictions of the inconsistent semantics which holds that these words shouldn't be appropriate for describing the same circumstance.

---

[8]We performed a similar pair of controls in Experiment 1B. We tested one control where the acceptability value was just the median acceptability of the first verb in the item, computed from acceptability judgments provided for that verb and that sentence frame in the norming study. We tested a second control where we computed an acceptability difference, but in this case we subtracted the median acceptability of the second verb from the first. In both cases, the controls did not impact the results of our hypothesis tests. The posterior estimates for all hypotheses were positive and the 95% credible intervals excluded zero.

Discussion

The results of this experiment support the pragmatic assumptions of our model. Participants are more inclined to accept an "in fact" cancellation statement for verb orderings that give rise to scalar implicatures than for verb orderings where the first verb implies the second verb. For example, participants generally find statements such as "The new technology affected the change, in fact it caused it." more acceptable than "The new technology caused the change, in fact it affected it." Moreover, participants' judgments of "in fact" statements that cancel implicatures generally skew to the acceptable side of the scale. The median acceptability is above the midpoint for all three of these orderings. The broad acceptability of the cancellation suggests that the use of a weaker verb does indeed implicate that the stronger verb isn't true.

Notably, the acceptability of statements that cancel implications is more mixed. The median acceptability for all three of these orderings hovers around the midpoint of the scale. This contrasts with the unacceptable statements in Experiment 1B (red striped bars in Figure 7) where participants' responses strongly skewed toward the unacceptable end of the scale. One way to make sense of this difference, is that the two unacceptable sentence types (e.g. "The new technology caused the change, but it didn't affect it." versus "The new technology caused the change, in fact it affected it") are unacceptable for different reasons. In the "but it didn't" constructions, the sentences that have verb orders where the stronger verb comes first are contradictory. It is not possible for something to "cause" but not "affect" an outcome, and so these types of sentences are seen as highly unacceptable. On the other hand, for the "in fact it" constructions, following the stronger verb with the weaker verb isn't contradictory, it is merely redundant. Stating the weaker verb in the second clause of the "in fact it" construction repeats something that was already implied by the first clause. A listener might find this strange, but not necessarily as unacceptable as an outright contradiction.

This speculation about different kinds of unacceptability colors the interpretation of

our model comparison. In outlining our predictions for the semantic overlap model for the caused → enabled ordering, we suggested that participants would rate this ordering as unacceptable. But we didn't distinguish between ratings of unacceptability due to contradiction and ratings of unacceptability due to unhelpful redundancy. In light of this distinction, the relative acceptability of caused → enabled ordering (compared to the same ordering in the "but it didn't" constructions) seems more understandable. The inconsistent semantics cannot avail itself of the same explanation. Under this semantics, the sentences should be unacceptable for the same reasons in both studies.

Moreover, the pattern of responses in enabled → caused ordering unambiguously supports the semantic overlap account over the inconsistent semantics. Participants clearly think these sentence frames are acceptable, indicating that 1) these two verbs are consistent, 2) "caused" is the stronger of the two, and 3) the use of "enabled" pragmatically implicates but does not imply that "caused" is not the case. These findings stand in contrast to the predictions of prior models, and highlight the benefit of developing an approach that can accommodate semantic overlap.

One caveat to this conclusion is that there is some ambiguity as to whether the Mental Model Theory proposes an overlapping or inconsistent semantics. The primary model presented for "enable" is the one we reproduce in Figure 1, and this is the model that applications of the theory use to generate predictions.[9] As we have shown, this model implies an inconsistent semantics. However both Goldvarg and Johnson-Laird (2001) and Khemlani et al. (2014) suggest another potential mental model for "enable" which includes all contingencies of cause and effect. This model is consistent with their model for "cause" and also less specific. The authors suggest that the more specific mental model for "enable" that they provide is pragmatically implicated, though they don't discuss how. Under this reading, the predictions of the Mental Model Theory would align with our own,

_____
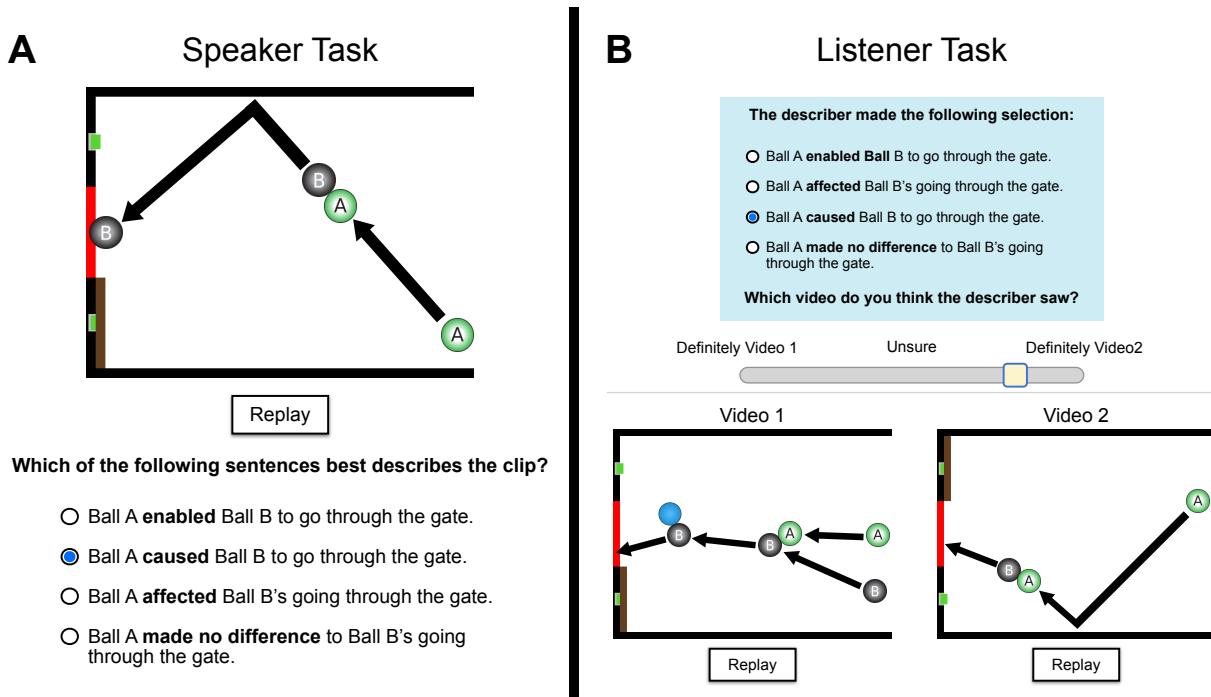
[9]It's worth noting again that the actual word Mental Model Theory provides a model for is "allow". As before we continue to elide the difference between these two words for now, but we will return to potential distinctions between "allow" and "enable" in the general discussion.

and in this case, one could see our model as elaborating the specific ways that mental representations, semantics, and pragmatics interact in people's use of causal language. These interactions are underspecified in the current form of the Mental Model Theory.

A final point to highlight is that "enabled" once again figures in the verb ordering with the most idiosyncratic participant responses (the caused → enabled ordering). This is perhaps still consistent with our theory given the considerations about redundancy versus contradiction, but it is nonetheless interesting to observe that in both Study 2 and Study 3 "enabled" interacts with one of the other verbs in a way that hints at added complexity. Plausibly there are additional semantic conditions that we haven't accounted for yet that may complicate the picture. For instance, Goldvarg and Johnson-Laird (2001) note that "enabled" seems to have the connotation that the result was intended, noting that it feels odd to say that "Shoddy work enabled the house to collapse." The implication of intentionality highlights the significance of agency, we might expect these connotations to be particularly strong when dealing wtih identifiable human agents (like the CEO). These additional complexities of "enable" require further consideration, which we will return to in the General Discussion.

Experiment 2: Speakers choosing what to say

Having validated the linguistic principles underlying our model, we now turn to quantitative assessments against participant behavior. We investigate how people choose what causal expression best describes what happened in dynamic physical scenarios. In our experiment, participants viewed physical scenarios like those illustrated in Figure 4, and chose from four causal expressions ("caused", "enabled", "affected", and "made no difference") the one that best describes the scenario. These physical scenarios allow us to quantitatively manipulate the different aspects of causation and see how this affects participants' use of the different causal expressions.

Figure 9. Sample trials from **A** the speaker task (Experiment 2), and **B** the listener task (Experiment 3). In the speaker task, participants chose one out of four utterances that best described what happened in the video clip. In the listener task, participants saw which one of the four utterances had been selected by a hypothetical speaker. Participants rated on a slider which one of two videos they thought the speaker had seen.

Methods

Participants.   We recruited 64 participants (*age*: M = 35, SD = 8, *gender*: 19 female, 43 male, 2 no response, *race*: 49 White, 6 Asian, 3 Black, 2 mixed race, 4 unclear or no response) online via Mechanical Turk using Psiturk (Gureckis et al., 2016). We excluded two participants from analyses who failed to select "made no difference" on an attention check video in which ball A lay still in a corner and thus clearly made no difference to the outcome. The experiment took 25 minutes on average (SD = 11), and participants were paid $3.67.

Stimuli.   We created 30 videos depicting physical scenarios, including the examples in Figure 3. All of the scenarios featured a ball labeled A, a ball labeled B, a red gate, a brown door, and two buttons controlling the door's movement (see Figure 4). Some of the scenarios also included a blue ball or a brown box. We constructed cases that varied the

causal aspects independently to create a range of causal scenarios emphasizing different aspects of causation (e.g. in some cases Ball A was a whether-, how-, and sufficient-cause. In others it was a how-cause and sufficient-cause, but not a whether-cause, etc...). We also included a number of classic causal scenarios that have influenced how researchers across disciplines have developed theories of causality. Schematics for all scenarios and corresponding aspect values are provided in Appendix D. Examples of classic causal scenarios include preemption scenarios (trials 6 and 23), causal chains (trials 4 and 14), Michottean launching (trial 13), and double prevention (trial 19).

Procedure. We screened for potential bots by asking a simple natural language question. Participants then received instructions about the task. We introduced the domain and the different objects in it, and had participants watch a video illustrating a scenario with all of the objects from the domain. Participants were told that they would view scenes like this one and then be asked to choose one out of the four descriptions that best captured the scene they viewed:

1. "Ball A **caused** ball B to go through the gate."

2. "Ball A **enabled** ball B to go through the gate."

3. "Ball A **affected** ball B's going through the gate."

4. "Ball A **made no difference** to ball B's going through the gate."

Participants then answered a comprehension check question. If they answered incorrectly, they were re-directed to the instructions to review them again. Once they successfully completed the comprehension check, participants advanced to the main task. Figure 9A displays a sample trial for the speaker task. Participants viewed the 30 test videos as well as one attention check video. The order of the videos was randomized. Below the video on each trial, we provided the prompt "Which of the following sentences best describes the clip?" followed by the four description options with radio buttons. The order of the first three descriptions was randomized between participants, but the

description with "made no difference" always came last. Participants had to view the video at least twice before making a selection. They were able to watch the video as many times as they liked, and chose to do so 2.2 times on average (SD = 0.6).

Analysis

Our model has four free parameters, θ in the causality module, which determines the amount of noise added to the objects' motions in the counterfactual simulations, σ and ν in the semantics module which determine, respectively, the softening for the additional features in the definition of "caused" (movement and uniqueness), and the softening in the definition of "made no difference", and λ in the pragmatics module which controls the speaker optimality. For a given value of these parameters, we can compute model predictions of the probability of selecting each causal expression on each trial. With these trial distributions we can then compute the likelihood of each participant response. We sum the log likelihood of all participant responses to assess likelihood of the data under our model at a given parameter setting. We fit the parameters using maximum likelihood optimization.

We compute aspect representations for each trial at several values of θ. We consider values ranging from 0.5 to 1.6 in increments of 0.1. Computing the aspect values is computationally expensive because it requires running many physics simulations, but the semantics and pragmatics computations are relatively cheap, and the remaining parameters can be fit using a fast optimization algorithm. We fit σ, ν, and λ using the L-BFGS-B algorithm (Byrd, Lu, Nocedal, & Zhu, 1995) implemented with Scipy's builtin optimizer (Virtanen et al., 2020). L-BFGS-B allows us to specify bounds for free parameters. We bound σ and ν at 0 and 1 as they represent the probability of responding with the given expression even if the condition they soften (movement, uniqueness, lack of how-cause) does not obtain. λ has a lower bound of zero but no upper bound. For each value of θ, we find the values of the remaining parameters that maximize the log likelihood of the data. We choose the model parameters that give the best fit across all these optimizations. We

found an optimal value 1.0 for θ, 0.65 for σ, 0.25 for ν, and 40.18 for λ.[10]

Alternative models. We compare our full model to two lesioned alternatives: a "No Pragmatics" model that removes the pragmatics component, and a "No Semantics and No Pragmatics" model that removes both components and computes a Bayesian ordinal regression instead, which directly maps from aspect values to utterance selections.

No Pragmatics This model removes the pragmatics component of the full model, and predicts selections based on a softmax function on the semantic values. While this model retains the semantic assumptions about the mapping between causal aspects and expressions, it does not consider how informative different utterances are. This model is analogous to a "literal speaker", who normalizes the semantic values across utterances instead of across scenarios (see Table 1b). We use the same parameter fitting procedure as in the full model with the same range for θ and parameter bounds. We found an optimal value of 1.0 for θ, 0.95 for σ, 1.0 for ν, and 2.54 for the temperature parameter of the softmax.[11]

No Semantics and No Pragmatics We fit a Bayesian ordinal regression with coefficients for each of the causal aspects, the movement feature, the uniqueness feature, random slopes and intercepts for each participant, and random intercepts for each trial. Because the predictors of the model are dependent on the noise parameter, θ, we fit one regression for each noise value sweeping across the same range of values we used for the other two models. We assumed the following ordering of expressions (from weakest to strongest): "made no difference", "affected", "enabled", "caused". While this ordering is consistent with our semantics, the regression assumes a linear additive mapping from

---

[10]We also considered whether to use one or two levels of recursive reasoning for our pragmatic speaker (three levels and above became computationally prohibitive in the grid search). The data was slightly more likely under the best-performing level-2 model than under the level-1 model. So we report the results for the level-2 pragmatic speaker here.
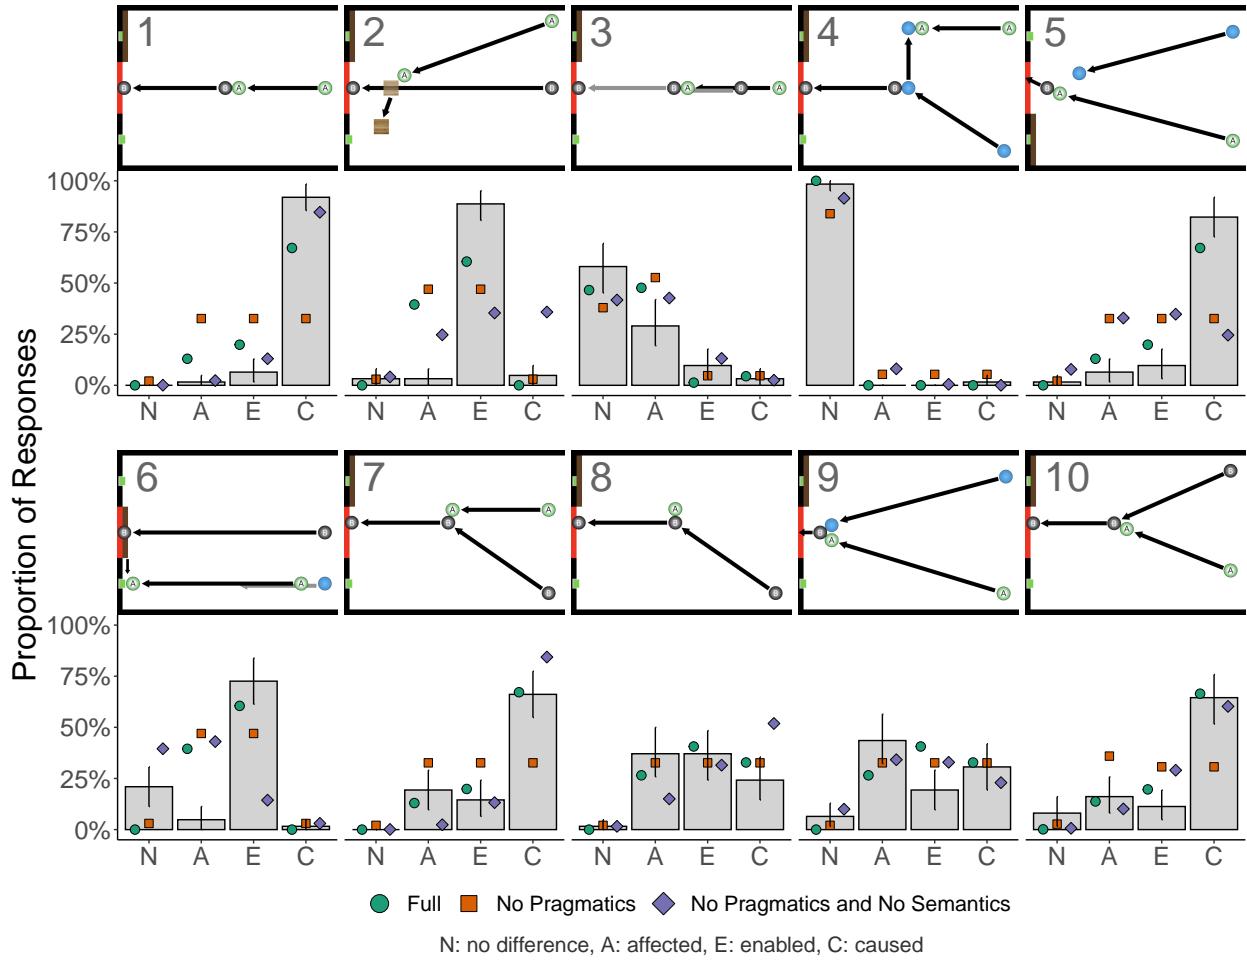
[11]For the "No Pragmatics" model, the value for ν is at the upper bound of the parameter range. We maintain the restriction on the range on this parameter given its interpretation as a probability. If we remove upper bound for this parameter, the "No Pragmatics" model improves slightly, but the overall pattern of results is unchanged.

causal aspects to expressions, rather than the logical semantic mapping of the full model. The regression model has eight parameters to estimate the fixed effects: three thresholds determining the boundaries between each of the causal expressions, and five coefficients determining the weight of each of the aspects of causation along with the movement and uniqueness feature. We found an optimal value of 1.0 for θ. Parameter estimates for this model are summarized in Table F1.

Semantics Analysis. In addition to the primary model comparison, we were interested to more closely examine our model semantics. The hierarchy of specificity illustrated in Experiment 1 constrains the space of possible definitions for our three causal expressions. But there are still many possible semantics that satisfy this constraint. How do we know that the particular semantics we specify above are appropriate? Our quantitative model allows us to evaluate different semantics and compare how well they explain the data we collected.

To perform this analysis, we first enumerated all the possible semantics for "affected", "enabled", and "caused" that were consistent with the hierarchy of specificity observed in Experiment 1. Our given model semantics satisfies this constraint, but so do other semantics. For example, if we define "caused" as $\mathcal{W} \wedge \mathcal{H} \wedge \mathcal{S}$, "enabled" as $\mathcal{W} \wedge \mathcal{S}$, and "affected" as $\mathcal{W} \vee \mathcal{S}$, this semantics would also exhibit the hierarchy of specificity. For this analysis we focused on the core aspect components of the definitions, whether-cause, how-cause, and sufficient-cause (excluding the movement and uniqueness features). We also fixed the definition of "made no difference" to the one we define in the model section above.

There are 258 semantics for "caused", "enabled", and "affected" that satisfy the specificity constraint. Once we had enumerated this entire set, we evaluated them each by fitting the model parameters to participant data from Experiment 2 and comparing the resulting log likelihood values. We fixed the θ parameter to 0.9 based on the value fitted in prior research on the CSM (Gerstenberg et al., 2021). This left two free parameters to fit for each semantics, ν the "made no difference" softener and λ the speaker optimality
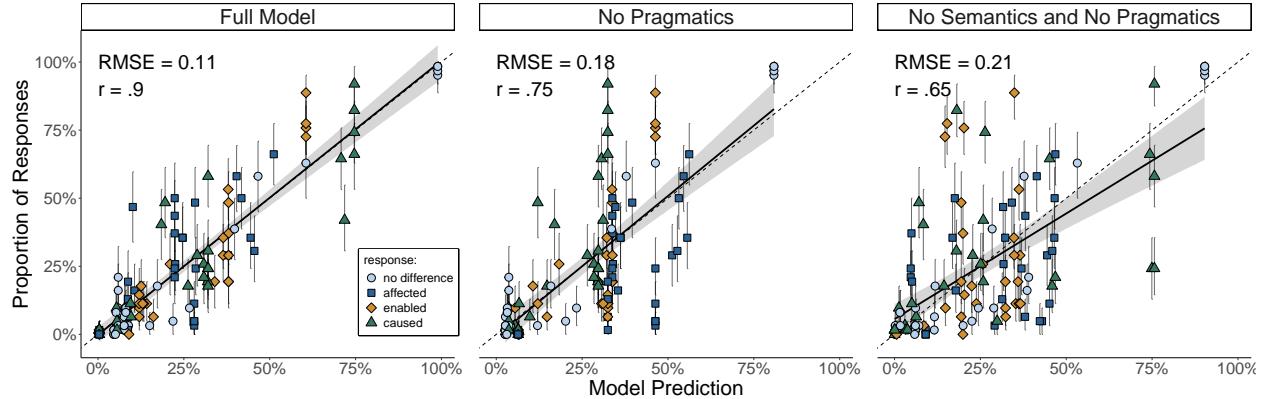
Figure 10. **Experiment 2**. Trial level model predictions for ten scenarios. In scenarios where participants strongly favor a particular utterance, the 'Full Model' captures this tendency, while the 'No Pragmatics' model assigns the same probability to all truthful utterances. Error bars represent bootstrapped 95% confidence intervals.

parameter.

Results

Figure 10 shows participants' selections for a subset of scenarios (Table E1 shows selections for all scenarios). In scenario 1, the classic Michottean case, participants strongly favor "caused", while in scenario 2 where ball A clears the box from ball B's path, participants strongly favor "enabled". In scenario 3, the modal response is "made no difference", though a substantial number of participants also selected "affected". In scenario 4, nearly every participant selected "made no difference". Scenarios 5 and 6

*Figure 11*. **Experiment 2**. Overall model performance for the "No Semantics and No Pragmatics" model, the "No Pragmatics" model, and the full model. Each point plots the model prediction for the probability of selecting a particular utterance on a given trial against the proportion of participants that selected that utterance on the same trial. The color and shape of the points indicates which utterance the point represents. There are thirty trials and four utterances per trial, so each panel contains 120 points. Overall, the full model performs best. The "No Pragmatics" model has a notable vertical line of responses around 33%. These reflect scenarios where multiple causal expressions are true, and so the "No Pragmatics" model ranks them all equally. The "No Semantics and No Pragmatics" model is more widely dispersed around the diagonal than the full model. Error bars and regression bands represent bootstrapped 95% confidence intervals.

illustrate preemption scenarios. In scenario 5, where ball A directly contacts ball B, participants favor "caused", while in scenario 6 where ball A hits a button opening the gate participants favor "enabled". Interestingly, in scenario 6, a substantial number of participants selected "made no difference". Scenarios 7 and 8 illustrate that movement matters, too. In scenario 7, where ball A is moving, participants strongly favor "caused", but in scenario 8, which is identical except ball A is stationary, participants are split between "affected", "enabled", and "caused". In scenario 9, where ball A and the blue ball together collide with ball B driving it through the gate, participants are also split between "affected", "enabled", and "caused" (the modal response is "affected"). In scenario 10, where there is some uncertainty as to whether ball B would have gone through without ball A being present, participants still strongly favor "caused".

Across these cases, the full model does the best job of capturing the data. In 7 out of 10 of the scenarios, the full model shows the closest match to the distribution of

participant responses. When participants strongly favor a particular response (e.g. scenario 1, 2, and 4), the full model's ability to select truthful and informative responses allows it to match this tendency. When participants show more variance in the expressions they select (e.g. scenario 3, 8, 9), the full model also matches this pattern. Both of the alternative models struggle with particular cases. The "No Pragmatics" model most clearly has issues with cases where participants favor "caused" (scenarios 1, 5, 8, 10). In these situations, "affected", "enabled", and "caused" are all true according to our semantics, so the "No Pragmatics" model rates them all equally. The "No Semantics and No Pragmatics" struggles with cases where participants favor "enabled" (scenario 2 and 6) as well as with scenario 5, likely because ball A is not a whether-cause in this case.

Figure 11 shows scatter plots of model predictions and aggregated participant responses for the full set of scenarios. The full model's predictions correlate best with participants' responses and show the lowest error, followed by the "No Pragmatics" model, and lastly the ordinal regression ("No Semantics and No Pragmatics"). In the "No Pragmatics" model, we can see a large column around 33% on the x-axis, representing the cases where there are multiple true utterances and the model weighs them all equally. In the "No Semantics and No Pragmatics" model the responses are in general more broadly dispersed than the responses for the full model.

To further assess model fit, we performed 100 split-half cross validation runs for each model, splitting the data by trials.[12] Table 8 presents the results. The full model performs the best, followed by the "No Pragmatics" model, and the "No Semantics and No Pragmatics" model.

The results of the semantics analysis are summarized in Appendix G. The most likely semantics for these data is in fact the semantics that we define above. This exploratory analysis further validates the particular semantics that we defined as the best semantics for

---

[12]Because we split the data by trials, we excluded random intercepts for trials for the ordinal regression in cross-validation.

Table 8

Experiment 2*. Speaker Experiment Split-Half Cross-Validation. The r column reports the median correlation coefficient on the test trials across the 100 cross-validation runs with 5% and 95% quantiles in brackets. The RMSE column reports the same for root mean square error. Δr reports the median difference in correlation coefficient between the Full model and the two alternative models, again with 5% and 95% quantiles in brackets. $\Delta RMSE$ reports the analogous difference in RMSE.*

Model	r	Δr	RMSE	ΔRMSE
Full Model	0.87 [0.80, 0.93]	–	0.13 [0.11, 0.16]	–
No Pragmatics	0.74 [0.63, 0.81]	0.14 [0.06, 0.21]	0.18 [0.15, 0.21]	0.05 [0.02, 0.08]
No Prag and No Sem	0.53 [0.27, 0.68]	0.35 [0.21, 0.58]	0.25 [0.20, 0.32]	0.12 [0.07, 0.17]

modeling the data in this task.

Discussion

In this experiment, we had participants select causal descriptions about what happened in various physical scenarios. This controlled setting allowed us to quantitatively compare model behavior to participant data. Overall, we see that each component of our model explains additional variance in participants' responses. This suggests that causal reasoning, semantics, and pragmatics are all important for understanding how participants choose to describe what happened.

Though in general, the full model does a good job of capturing participant responses, there are some situations where its predictions are off. Notably, in cases like scenario 2 and scenario 6 where "enabled" is the dominant response, the full model over-weights the probability of "affected", even though almost no participant used this utterance for these scenarios. One potential explanation is that our model pragmatics does not accurately capture the informativeness of using "enabled" when "affected" is also true. The tendency to favor the more informative utterance is determined by how often that utterance is true relative to the less informative utterance. The fewer worlds "enabled" is true of relative to "affected", the greater its informativity and the more our pragmatics model will favor it when both words are true. The pattern we see here could be explained by the fact that participants think "enabled" is substantially more informative than our model does.

Another notable detail about scenario 6 is that a number of participants (around 20%) select "made no difference". Scenario 6 is a situation of "preempted enablement": ball A presses the button that moves the brown door out of the way allowing ball B to pass through the gate, but even if ball A hadn't pressed the button, the blue ball would have done so. Interestingly, for the other situation where "enabled" is the dominant response (scenario 2) and the other case of preemption (scenario 5), almost no participants respond "made no difference".

Finally, it's interesting that participants almost never selected the expression "enabled" in scenario 3. Most participants said that ball A either "made no difference" to the outcome, or that it "affected" ball B's going through the gate – and this is predicted by the CSM. In contrast, the force dynamics model (Wolff, 2007) predicts that participants should say "enabled" here. In fact, the force configuration maps exactly onto the paradigmatic case for "enabled" as shown in Figure 1. One possible way of resolving this discrepancy is to assume that in many instances, there is some degree of uncertainty about whether the patient would have achieved the endstate by itself. For example, if there is friction then it's possible that the patient would not have reached the endstate without the help of the agent. However, notice that this kind of reasoning is counterfactual and thereby outside of the purview of what the force dynamics model considers (although see Wolff et al., 2010, for an extension of the force dynamics model that incorporates some counterfactual machinery).

Experiment 3: Listeners inferring what happened

The previous experiment focused on the speaker side of communication. A speaker who saw what happened chooses what description to use. The RSA framework allows us to easily pivot our model to make predictions in the listener setting, too. In this third experiment, we examine what a listener can infer about the scenario given a causal description of what happened.

Methods

Stimuli.   Figure 9B shows a sample trial for the listener task. Each trial consisted of a description (one of the four utterances from Experiment 1) and a pair of video scenarios. The scenarios were selected from the same set of videos used in Experiment 2. To select what pairs to show on a trial, we considered every possible video pair and evaluated their relative probability given each utterance under the pragmatic listener. For each utterance, we selected nine video pairs, varying the absolute difference in relative probabilities within each utterance. On some trials the model strongly preferred one video over the other, but on other trials, the model had a weak preference or no preference at all. In all trials, the utterance that was attributed to the speaker was participants' modal response in Experiment 2 for at least one of the scenarios shown in the trial.

Participants.   We recruited 71 participants online via Mechanical Turk using Psiturk (*age*: M = 37, SD = 10, *gender*: 19 female, 51 male, 1 non-binary, *race*: 40 White, 21 Black, 8 Asian, 2 unclear or no response). We removed 21 participants who failed to pass an attention check, leaving us with 50 participants for analysis.[13] The experiment took on average 31 minutes (SD = 13), and each participant was paid $5.50.

Procedure.   Participants first received instructions about the physical setting and the different objects in it. We then instructed participants on the speaker task from Experiment 2 and had them complete a comprehension check. After they passed the check, participants completed a short "training session" where they performed the speaker task for four trials. We then instructed participants on the listener task. Participants were told that rather than selecting a description for a video, they would now see a description that someone else chose. Their task was to indicate which of two scenarios they thought the hypothetical "describer" had seen, based on the chosen description. Participants completed another comprehension check for this new task and then proceeded to the main phase of

---

[13]Attention check attrition was higher in this experiment. This could be in part because the check included two trials and was thus more stringent.

the experiment. Participants who failed either comprehension check were sent back to the instructions for the corresponding experiment section. Participants needed to pass all comprehension checks to proceed to the main phase of the experiment.

On each trial, participants saw which out of the four causal expression the describer had selected, indicated by a highlighted radio button next to the chosen expression (see Figure 9B). Participants were asked to answer the question: "Which video do you think the describer saw?" Below the prompt were two videos labeled "Video 1" on the left and "Video 2" on the right. Participants had to first watch the video on the left and then the video on the right. After they had watched each video once, a sliding scale appeared above the videos. The endpoints of the slider were labeled "Definitely Video 1" on the left, "Definitely Video 2" on the right, and the midpoint was labeled "Unsure". The scale ranged from -50 to 50, though there were no numeric values visible to participants. To avoid anchoring effects, the slider handle was initially invisible and only appeared upon clicking on the slider. After a participant had viewed each video once and indicated a judgment on the slider, they could proceed to the next trial. Participants could replay videos as many times as they wanted (video play count M = 1.13, SD = 0.41).

Participants provided judgments on 36 trials and 2 attention checks. In one attention check, the describer's chosen expression was "caused", and one video illustrated a strong causal role of ball A while the other lacked any causal connection between ball A and ball B. In the second attention check, the chosen expression was "made no difference", and again the comparison in the two scenarios was chosen to be maximally salient. Participants who put the slider on the opposite side of the correct video on at least one of the two attention checks were excluded from the analysis. The order of the trials and the position of the videos on each trial (left or right) was randomized between participants.

Analysis

We rescaled participant responses to lie on the interval $[0, 1]$, where a zero response represents a judgment favoring the left scenario, and 1 represents a judgment favoring the
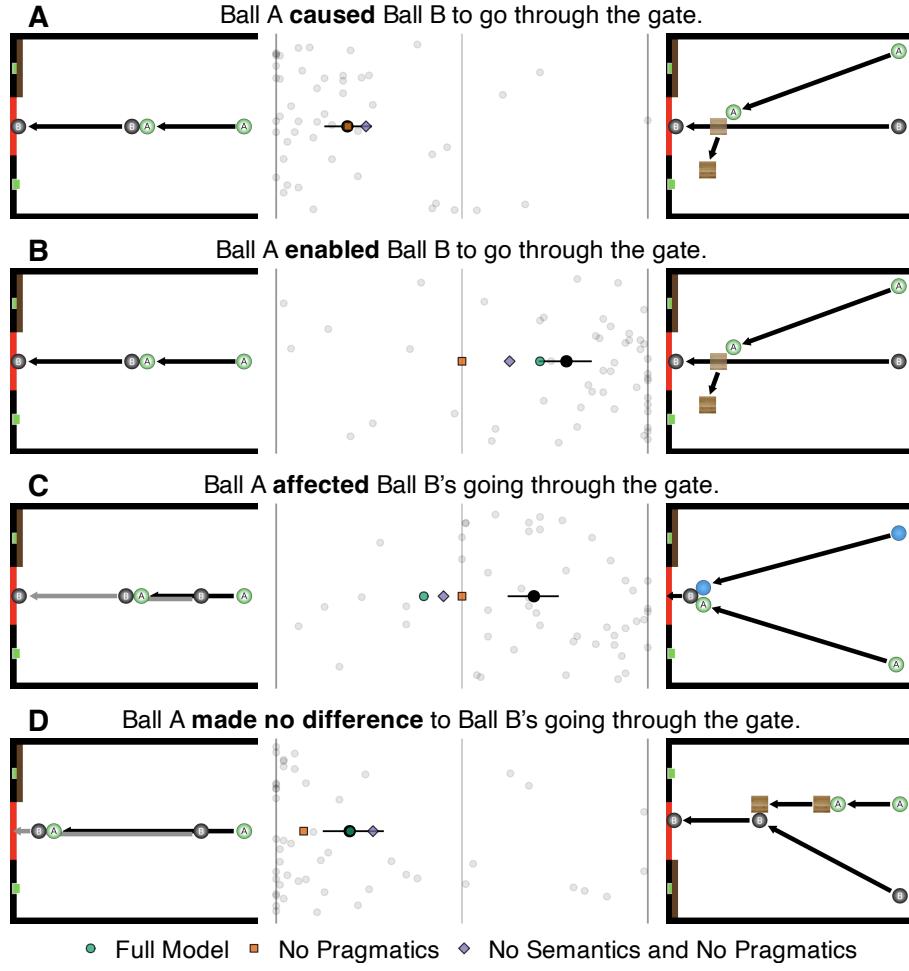
right scenario. We computed the mean participant response for each trial.

With the model parameters that we fit in the speaker task, we used a level-1 pragmatic listener to compute a distribution on scenarios given each utterance. Each trial in our experiment consists of a given utterance and a pair of scenarios. To make a prediction for a particular trial, we took the probability of each scenario under the given utterance and applied a softmax with temperature parameter β. We fitted β to minimize the squared error between model predictions and participant means.

Alternative Models.   We included two alternative models analogous to the alternatives from Experiment 2. For the "No Pragmatics" model, we used a literal listener instead of a pragmatic listener to compute the distribution over scenarios given each utterance. For the "No Semantics and No Pragmatics" model, we took the predictions from the best-fitting ordinal regression. We fitted the predictions for both models using a softmax function (with separate β parameters for each model).

Results
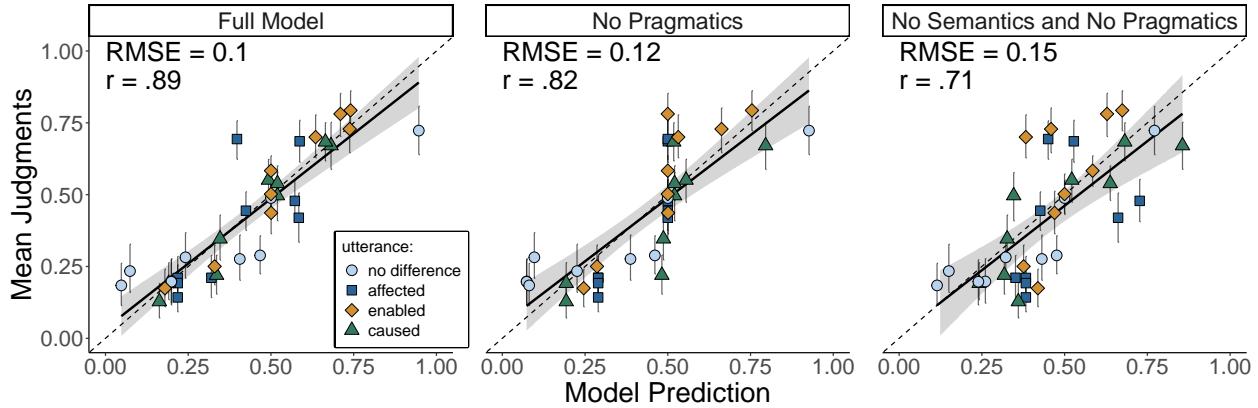
Figure 12 shows participant responses for four trials (Table H1 shows the responses for all trials). In Figure 12A, the speaker said that "Ball A **caused** ball B to go through the gate." In the left scenario, ball A knocks into ball B and ball B goes through the gate. In the right scenario, ball A knocks into a box pushing it out of ball B's path and allowing ball B to go through the gate. Participants judged that it was more likely that the speaker had referred to the scenario on the left, and all three models capture this preference. In Figure 12B, participants saw the same two scenarios, but this time the speaker said that "Ball A **enabled** ball B to go through the gate." Here, participants strongly favored the scenario on the right. The full model matches participants' responses more closely than the alternatives. Notably, the "No Pragmatics" model predicts that each scenario is equally likely because the utterance is true in both scenarios. The full model, like participants, draws the pragmatic inference that the speaker would have used the stronger utterance "caused" had they seen the scenario on the left, so they must have seen the scenario on the

Figure 12. **Experiment 3**. Participant responses and model predictions for a selection of cases. Videos are illustrated on either end with the trial utterance listed above. Dark black dot represents participant mean selection with bootstrapped 95% confidence intervals. Light dots represent individual participant judgments. Colored shapes represent model predictions. **A)** All models do a good job capturing participant responses for the expression "caused". **B)** For the same scenarios as in **A** when the expression is "enabled", the "No Pragmatics" model fails to capture participant responses. Pragmatic inference allows the full model to infer the speaker intends to communicate "enabled" *but not* "caused" like participants do. **C)** All models fail to capture participants' response tendency. Participants judge that "affected" better describes the scenario on the right, while the "Full Model" predicts that participants would prefer the scenario on the left. **D)** All models do a good job capturing participants' preference for the left scenario given that "made no difference" was selected.

right.

In Figure 12C, the speaker said that "Ball A **affected** ball B's going through the

Figure 13. **Experiment 3**. Overall performance for the full model and alternatives in the listener experiment. Each point represents the model prediction for a particular trial compared against the mean participant response on that trial. The utterance type is indicated by the color and shape of the point. Error bars and regression bands represent bootstrapped 95% confidence intervals.

gate." On the left, participants saw a case where ball B is headed toward the gate on its own, ball A comes up behind it and pushes it along. In this scenario, ball A is only a how-cause (but not a whether-cause or a sufficient-cause). On the right, they saw a case where ball A and the blue ball collide with ball B simultaneously driving it through the gate. In this scenario, ball A is a how-cause and a sufficient-cause. According to our semantics, "affected" is true for both scenarios, but "enabled" and "caused" are also true of the scenario on the right. Interestingly, participants strongly favored the scenario on the right. All of the models failed to capture this tendency. In this case, the full model erred the most because it uses pragmatic reasoning to infer scenario on the left (because the speaker could have used a stronger expression had they wanted to refer to the video on the right). In Figure 12D, the speaker said that "Ball A **made no difference** to ball B's going through the gate." On the left, participants saw a scenario where ball A pushes ball B along after it is already headed to the gate. On the right, participants saw a scenario where ball A knocks into a box which then knocks into ball B and re-directs it through the gate. Participants strongly favored the scenario on the left here, and all models capture this pattern.

Table 9

Experiment 3*. Listener Experiment Split-Half Cross-Validation. The r column reports the median correlation coefficient on the test trials across the 100 cross-validation runs with 5% and 95% quantiles in brackets. The RMSE column reports the same for root mean square error. Δr reports the median difference in correlation coefficient between the full model and the two alternative models, again with 5% and 95% quantiles in brackets. $\Delta RMSE$ reports the analogous difference in RMSE.*

Model	r	Δr	RMSE	ΔRMSE
Full Model	0.91 [0.84, 0.94]	–	0.10 [0.08, 0.12]	–
No Pragmatics	0.83 [0.75, 0.89]	0.07 [0.01, 0.15]	0.13 [0.10, 0.16]	0.03 [0.0, 0.06]
No Prag and No Sem	0.72 [0.60, 0.81]	0.19 [0.09, 0.30]	0.16 [0.14, 0.18]	0.06 [0.03, 0.08]

Figure 13 illustrates overall model performance for the full model and the two alternatives. The pattern of results is similar to the speaker experiment. The full model does the best job predicting participant responses, followed by the "No Pragmatics" model, and finally the "No Semantics and No Pragmatics Model". For the "No Pragmatics" model, there is again a column of responses near the middle. This is the case because the "No Pragmatics" can't distinguish between scenarios on trials in which the given utterance is true of both scenarios.

As with the speaker experiment, we ran 100 split-half cross-validations to compare model fits. The results are summarized in Table 9. Again we see that the full model outperforms the alternatives. This time, the difference in performance between the full model and the "No Pragmatics" model is smaller than in Experiment 2.

Discussion

In this experiment we demonstrate that our model explains patterns of participant behavior in the listener setting. In terms of model performance, the pattern of results is very similar to what we saw in the speaker experiment (Experiment 2). The full model performs better than both of the alternatives in the cross-validation. This time, the performance difference between the full model and the "No Pragmatics" was smaller. However, the full model holds a qualitative advantage over the "No Pragmatics" model in that it correctly predicts participant responses on trials that involve scalar implicatures. In

trials like the one presented in Figure 12B, the full model infers that because the speaker used a weaker utterance, they intended to communicate that the ball "enabled" but didn't "cause" the outcome. This behavior is consistent with the semantic and pragmatic assumptions of our model which we validated in Experiment 1.

Certain cases still pose a challenge for our model. This is most clearly apparent in the trial depicted in Figure 12C. Though all the models fail to capture the general pattern of participant responses in this case, the full model is especially off. Because the full model believes stronger utterances are true of the scenario on the right, it favors the scenario on the left for which "affected" (and to some extent "made no difference") are the only true utterances. Interestingly, in the speaker experiment, a plurality of participants favored "affected" for the scenario on the right (see scenario 9 in Figure 10).

For the listener task, we can also fit an empirical model based on participant responses in the speaker task. For each trial, this model takes the proportion of participants that responded with the given utterance on the two scenarios in the speaker experiment and then normalizes these two values using a softmax function (with a fitted temperature parameter β). For example, in Figure 12A, the model would take the proportion of participants from Experiment 2 who responded "caused" in the scenario on the left and on the right, and then run this pair of values through a fitted softmax. This empirical model captures participants' inferences in the listener experiment very well with $r = .95, \text{RMSE} = 0.07$. The fact that the empirical model performs so well suggests that the listener task doesn't introduce many additional factors that go beyond what participants need to do in the speaker task. As our model suggests, a listener can infer what happened simply by considering the extent to which they would have used a given expression in the different situations (cf. Amemiya, Heyman, & Gerstenberg, 2024; Kirfel, Icard, & Gerstenberg, 2022).

General Discussion

Causality permeates our everyday language in ways both subtle and pronounced. In this paper, we developed the counterfactual simulation model of people's use of different causal expressions including "caused", "enabled", "affected", and "made no difference". The model draws insights from philosophy, linguistics, and psychology. It consists of three modules: a causal knowledge module, a semantics module, and a pragmatics module. The *causal knowledge module* computes a representation of different causal aspects of a scene by simulating the consequences of different counterfactual interventions. The *semantics module* defines different causal expressions in terms of logical combinations of these causal aspects. The *pragmatics module* then computes which utterance to use (or scenario to infer) based on principles of rational communication. Together these three components offer an account of how people choose and interpret causal expressions.

We tested our model in a series of experiments. In an initial set of psycholinguistic studies, we validated the model's semantics and provided evidence that people draw pragmatic inferences as predicted by our model. After running a norming study (Experiment 1A), we asked participants to rate the acceptability of a series of sentences that tested our model semantics (Experiment 1B) and pragmatics (Experiment 1C). We found that the qualitative pattern of participants' judgments was consistent with our model's predictions. Furthermore, we demonstrated that these results support a semantic account like ours which posits that the meanings of the causal expressions "caused" and "enabled" are overlapping, and contrasts with prior work suggesting that these are inconsistent causal concepts.

We followed up this qualitative validation and model comparison with two quantitative tests. In Experiment 2, participants took the role of a speaker. They viewed videos of physical interactions and chose what sentence best described what happened. In Experiment 3, participants took the role of a listener. Their task was to infer which of two videos a speaker had seen based on the causal expression they had selected. Our model

captured participants' judgments well in both these tasks. To assess whether each component of our model was critical, we compared the full model with two lesioned alternative models, removing either only the pragmatics module, or both the pragmatics and the semantics modules. In both experiments, we found that the full model outperformed the alternatives in cross-validation, suggesting that all three components of the model are important for understanding how people perform these tasks.

Our model presents a novel contribution in that it explicitly captures the interaction between causal concepts and linguistic communication. Though some linguists have previously observed that pragmatic considerations impact the meaning of causal expressions (e.g. McCawley, 1978), these considerations have yet to be fully taken up by psychologists. Prior accounts have proposed direct mappings between mental representations and causal expressions (Goldvarg & Johnson-Laird, 2001; Sloman et al., 2009; Wolff, 2007). Though some observe that pragmatics plays a role in the meaning that is ultimately inferred (Goldvarg & Johnson-Laird, 2001; Khemlani et al., 2014), they don't elaborate how. Articulating the relationship between causal representations, semantic meaning, and pragmatic inference brings to light subtleties in the meanings of causal expressions that have been clouded by previous approaches. For example, prior work had suggested – implicitly or explicitly – that the meanings of "caused" and "enabled" are inconsistent (Goldvarg & Johnson-Laird, 2001; Sloman et al., 2009; Wolff, 2007). In contrast, we provide empirical evidence that the meanings of these words overlap, and that people resolve semantic ambiguities through pragmatic inference. Our semantics builds on and extends these prior accounts. Like Goldvarg and Johnson-Laird (2001) and Sloman et al. (2009), we highlight the importance of necessity and sufficiency through the notions of whether-causation and sufficient-causation. Like Wolff (2007), we incorporate the importance of process through the notion of how-causation. But by showing how these underlying concepts are combined in the meanings of these causal expressions and then differentiated through pragmatic inference, we illustrate a more nuanced psychological

process that can help us to make sense of the new findings we present here.

Our work establishes a framework for modeling the relationship between people's causal representations of the world and their causal language. The implementation of this model requires various assumptions. In the remainder, we will consider some limitations of our implementation that suggest potential avenues for future research. We structure our discussion by going through the three different modules one by one: causal knowledge, semantics, and pragmatics.

Causal Knowledge

The causal knowledge module is based on the counterfactual simulation model (CSM) (Gerstenberg et al., 2021). The CSM was developed to model causal judgments in intuitive physical settings, and that was the domain we focused on here for quantitative modeling. In some ways the physical domain is quite simple, but it still allows us to capture a wide variety of causal scenarios that have driven discussion about causality across disciplines (e.g. causal chains, double prevention, preemption, etc.). Still, people use causal language across a wide variety of domains, and a more comprehensive model would capture these different kinds of causal interactions.

New questions of implementation arise when generalizing the tools of the CSM to new settings. As a starting point, the CSM requires a generative model which supports counterfactual simulation in the domain of interest. Recent work has begun to implement these models for social causal interactions (Sosa, Ullman, Tenenbaum, Gershman, & Gerstenberg, 2021; Wu, Sridhar, & Gerstenberg, 2022; Wu et al., 2023). Rather than relying on noisy physics simulators, these approaches use models of planning and rational action (Jara-Ettinger, Gweon, Schulz, & Tenenbaum, 2016; Ullman et al., 2009). They describe how agents select actions by weighing costs and benefits and how they update their beliefs about the world. These models also support inferences about the latent beliefs and desires that give rise to an observed set of actions (Baker, Jara-Ettinger, Saxe, & Tenenbaum, 2017; Baker, Saxe, & Tenenbaum, 2009).

With a new set of generative models, additional questions arise about how to compute the different causal aspects. When testing for whether-causation in the physical domain, removing the candidate cause is a straightforward way to assess counterfactual necessity. In the social domain, the appropriate counterfactual operation is less obvious. One possible contrast is removing the agent altogether, but this might not be the standard counterfactual that people rely on when assessing social causal roles. For instance, in the law, the reasonable person test (Jackson, 2013) is a counterfactual test that prompts jurors and judges to compare the behavior of a defendant to an imagined reasonable person (cf. Gerstenberg et al., 2018; Wu & Gerstenberg, 2024). The agent isn't removed from the setting; rather their mental states are manipulated such that they meet a standard of reasonableness. Generative models of social cognition can support these kinds of counterfactuals, though defining precisely what kind of operations are required to make someone reasonable requires further research at the intersection of social cognition and law (Lagnado, Fenton, & Neil, 2013; Summers, 2018; Tobia, 2018).

Though the physical and social world are two of the most pervasive domains of causal thinking, people's causal thought ranges further still, supporting judgments about topics as diverse as human biology ("Eating the ice cream gave me stomach ache."), weather ("The heavy rains caused the flowers to bloom."), and political-economic events ("The war in Ukraine made gas prices go up."). In principle, the CSM could generalize to model causal thinking in such domains by operating over increasingly abstract generative models (Beckers & Halpern, 2019; Ho et al., 2022; Shin & Gerstenberg, 2023). Critically, these models need not maintain a detailed representation of the processes that support causal reasoning in the physical and social domains. Laypeople making judgments about the causal relationships in complex phenomena likely operate with relatively simplified intuitive models about the mechanisms that underlie those relationships (Rozenblit & Keil, 2002). Even experts in these fields need to make various simplifying assumptions in order to model complex phenomena at the level of, for instance, economic systems. As long as

the model supports counterfactual simulation, it can be used to reason about cause and effect (Tavares, Koppel, Zhang, Das, & Solar-Lezama, 2021).

Semantics

In our model, we propose a semantics for three periphrastic causatives: "caused", "enabled", and "affected". According to our semantics, these expressions exist on a hierarchy of specificity, where "caused" is the most specific, "enabled" is in the middle, and "affected" is the least specific. In Experiment 1, we showed that participants broadly affirm this hierarchy for causal statements describing a wide range of causal phenomena. In Experiment 2, we further demonstrated that the particular semantics we defined provides the best quantitative fit to the data that we collected.

Still there are limits to the generality of these claims. Examining our norming study (Experiment 1A), one may object that the existence of sentence frames where the stronger utterance is acceptable but the weaker utterances are not acceptable provides a counterexample to our semantics (e.g. bottom row of Figure 5). We exclude these types of cases from our subsequent tests, but these are precisely cases that would contradict our tests if we had included them. In fact, Wolff (2007, p. 84) considers a similar sentence frame where "caused" is acceptable and "enabled" is not in coming to the conclusion that the two verbs are inconsistent:

1. A cold wind caused him to close the window.

2. ?A cold wind enabled him to close the window.

Cases like these do imply limits on the generality of our semantics. Though we can't discount the evidence from Experiment 1 suggesting semantic overlap is present in a wide variety of causal domains, we also must acknowledge that there seem to be cases where it is not the case. Understanding the bounds of where these semantics do and don't apply is an important direction for future work.

An important part of developing this understanding is further unpacking the slippery

semantics of "enabled". Notably, in Experiment 1, both situations where participant responses deviated from our model predictions involved this expression (the affected → enabled ordering in Experiment 1B and, arguably, the caused → enabled ordering in Experiment 1C). It appears that there are additional semantic considerations for this word in particular that we have yet to account for. The Oxford English Dictionary defines "enable" like so: "give (someone or something) the authority or means to do something" or to "make possible". There does seem to be a counterfactual notion of whether-causation here, but there are other aspects, too. The provision of "authority" seems to suggest social relations perhaps of unequal power. Indeed, "enable" seems to have additional connotations in social settings. In these contexts, it seems strange to use "enabled" if the person being enabled didn't desire the outcome. For example, saying that "The heavy traffic enabled Caroline to miss her flight" sounds odd, unless we know that Caroline actually wanted to miss it and was headed to the airport reluctantly. On the other hand, it seems fine to say that "The heavy traffic caused Caroline to miss her flight" regardless of whether she wanted to or not. Consistency with the desire of an enabled person is one potential semantic restriction of "enabled". In another definition suggested by Cao, Geiger, Kreiss, Icard, and Gerstenberg (2023), "enabled" is defined as "causing another agent to be able to achieve some outcome". This aligns well with the "making it possible" sense of enabled. Cao et al. (2023) also work in a social setting, and interestingly, their definition doesn't actually require that the outcome itself came to fruition. Their experiments seem to suggest that a number of participants are willing to endorse statements with "enabled" even when the outcome didn't occur.

In addition to better understanding the meanings of the causal expressions we currently model, future work should expand the set of expressions. Beyond "caused", "enabled", and "affected", linguists and psychologists have analyzed a wide variety of periphrastic causatives. For example, Nadathur and Lauer (2020) explore the semantics of "make", highlighting the importance of sufficiency in its meaning. Sufficiency is already

implemented in the CSM, and it would be straightforward to provide a definition in the causal physical setting and then test that definition alongside our other expressions. Another interesting set of causatives to consider are what Wolff and Song (2003) call "enable-type" verbs. These are verbs like "allow" and "let" which have similar meanings to "enable". In this work we've treated "allow" and "enable" as roughly equivalent, as do Wolff (2007) and Sloman et al. (2009), though Goldvarg and Johnson-Laird (2001) note certain nuanced differences. One interesting observation is that, in the social domain, "allow" may not carry with it the same semantic constraint noted a moment ago for "enable" that the patient desires the outcome. By our own intuition, it seems acceptable to say that "the teacher allowed their student to fail the test", while it feels stranger to say that "the teacher enabled their student to fail the test". Still these intuitions require further validation among a broader population of speakers. Our framework provides the potential to make explicit hypotheses about the differences (and similarities) in meanings among these verbs and then test those hypotheses quantitatively.

Our model focuses on analyzing the meanings of individual words, but ultimately when testing participants we ask them to provide judgments about entire sentences. We equate the meanings of the word and the sentence, but this is a simplification. Sentential context can impact the meanings of the individual words that make up the sentence. This is perhaps most clear in the case of "affected". We define "affected" as the disjunction of all the aspects, and the phrasing that we use is meant to allow for the most inclusive possible meaning ("Ball A affected Ball B's going through the gate"). Changing the surrounding sentence could impact how people interpret the word. If we were to say instead that "Ball A affected **whether** Ball B went through the gate" or "Ball A affected **how** Ball B went through the gate", participants might interpret "affected" as denoting the particular aspect of causation being highlighted. Understanding how the meanings of these individual words compose with others to give rise to the entire sentence meaning is an important part of filling out the linguistic picture.

An important limitation of our work is that, so far, the model only applies to a set of English causatives. In the future, we would like to expand this approach to other languages as well. Cross-linguistic research has revealed interesting similarities but also important differences in the ways that causation is expressed and understood across languages (Beller, Song, & Bender, 2009; Bender & Beller, 2011; Comrie, 1976; Haspelmath, 2016; Wolff, Jeon, & Li, 2009; Wolff et al., 2005). The psychological representations that underlie the use of a word in one language and its translated equivalent in another might not be the same. For example, Klettke and Wolff (2003) find that English and German participants differ in their tendencies to describe the same scenarios using "cause" (German "verursachen") or "enable" ("ermöglichen"). Our model provides a framework to explore the psychological underpinnings of these differences.

Pragmatics

In our model, we used the Rational Speech Acts (RSA) framework (Degen, 2023; Frank & Goodman, 2012; Goodman & Frank, 2016) to model participants' use of scalar implicature. In the literature on scalar implicature, linguists often distinguish between lexicalized scales and ad hoc scales (Degen, 2015; Hirschberg, 1985). Whereas in a lexicalized scale, the use of a weaker utterance always invites the stronger utterance as an alternative (as is often suggested for "some" and "all"), in an ad hoc scale, additional context may be required to make the comparison salient. It is difficult to assess whether the scale we've identified in this work is lexicalized or ad hoc. The setup in our experiments makes the relevant contrast very clear. While this does not rule out the possibility that people automatically consider "cause" as an alternative when hearing "enable" in natural speech, we would need different methods to test what alternatives naturally come to people's minds.

The existence of ad hoc scales, however, highlights the role of context and comparison in pragmatic speech, and raises questions about how the contrast set might impact our analysis. To some extent, the effects we've observed are shaped by the set of alternative

utterances that we had participants consider. Naturally, we might wonder what would happen if we changed the set of alternatives. For example, would people use "caused" in our scenarios differently when "made" was an alternative? Together with questions about the semantics of other causative constructions come questions about the pragmatics. If we define additional causative constructions, how does pragmatic inference drive the choice of one or another given a particular context? What if we contrast lexical causatives with periphrastic ones? The linguistics literature has explored extensively the conditions under which it is appropriate to use a lexical causative or a periphrastic alternative (Fodor, 1970; Katz, 1970; McCawley, 1978; Shibatani, 1976), though sometimes these examinations are performed by individual linguists on a small number of examples. Our model provides a new framework for examining these questions, allowing us to formulate explicit hypotheses about the semantics and pragmatics of causal verbs.

These directions for future study highlight the inter-relatedness of the semantic and pragmatic modules of our model. Though in our work here, these modules are neatly separated, studies in linguistics reveal that the assumed separation between these capacities is often not so clear-cut (Börjesson, 2014). Addressing this observation, Bergen, Levy, and Goodman (2016) developed an RSA framework where pragmatic inference extends to the semantic content of the expressions themselves. Rather than assuming the semantic content of utterances is fixed across contexts, this approach assumes that interlocutors maintain a level of lexical uncertainty over the semantics of different utterances. In addition to inferring which utterance a speaker would choose in context, a listener also infers what exactly the speaker means by an utterance (see also Potts, Lassiter, Levy, & Frank, 2015). Extending our model in this way could expand its explanatory power and help address problem cases like the one from Wolff (2007) which seems to suggest additional semantic properties of "enabled" that we have not accounted for.

Conclusion

Causation is complex and multi-faceted. There are many ways in which one event can make a difference to another, and our language provides us with a limited set of expressions to communicate what happened. Did Charlie cause Phil's success, or merely enable it? To capture how people use and understand various causal expressions, we developed and tested the *counterfactual simulation model of causal language*. The model assumes that people form a causal representation of what happened by paying attention to the way in which the candidate cause brought about the outcome. A cause can make a difference to whether and how an outcome happened, and this is revealed through simulations of what would have happened in relevant counterfactual situations. The meaning of different causal expressions, such as "caused", "enabled", "affected", and "made no difference" is then defined in terms of these aspects of causation, and by incorporating pragmatic principles of rational communication, the model accounts for how speakers choose what words to use, and how listeners infer what happened. By drawing on insights from philosophy, linguistics, and psychology, the work presented here brings us one step closer to understanding how people communicate about causality.

References

Amemiya, J., Heyman, G. D., & Gerstenberg, T. (2024). Children use disagreement to infer what happened. *Cognition*.

Aronson, J. L. (1971). On the grammar of 'cause'. *Synthese*, *22*(3), 414–430.

Baglini, R., & Siegal, E. A. B.-A. (2021). Modelling linguistic causation. *manuscript, Aarhus University and Hebrew University of Jerusalem*.

Baker, C. L., Jara-Ettinger, J., Saxe, R., & Tenenbaum, J. B. (2017). Rational quantitative attribution of beliefs, desires and percepts in human mentalizing. *Nature Human Behaviour*, *1*(4), 0064.

Baker, C. L., Saxe, R., & Tenenbaum, J. B. (2009). Action understanding as inverse planning. *Cognition*, *113*(3), 329–349.

Battaglia, P. W., Hamrick, J. B., & Tenenbaum, J. B. (2013). Simulation as an engine of physical scene understanding. *Proceedings of the National Academy of Sciences*, *110*(45), 18327–18332.

Beckers, S. (2021). Causal sufficiency and actual causation. *Journal of Philosophical Logic*, *50*(6), 1341–1374.

Beckers, S., & Halpern, J. Y. (2019). Abstracting causal models. In *Proceedings of the aaai conference on artificial intelligence* (Vol. 33, pp. 2678–2685).

Beller, A., Bennett, E., & Gerstenberg, T. (2020). The language of causation. In *Proceedings of the 42nd Annual Conference of the Cognitive Science Society*.

Beller, S., Song, J., & Bender, A. (2009). Weighing up physical causes: Effects of culture, linguistic cues and content. *Journal of Cognition and Culture*, *9*(3), 347–365.

Bender, A., & Beller, S. (2011). Causal asymmetry across cultures: assigning causal roles in symmetric physical settings. *Frontiers in psychology*, *2*, 231.

Bergen, L., Levy, R., & Goodman, N. (2016). Pragmatic reasoning through semantic inference. *Semantics and Pragmatics*, *9*, 1–83.

Börjesson, K. (2014). The semantics-pragmatics controversy. In *The semantics-pragmatics*

controversy. de Gruyter.

Bunzl, M. (1980). Causal preemption and counterfactuals. *Philosophical Studies: An International Journal for Philosophy in the Analytic Tradition*, *37*(2), 115–124.

Bürkner, P.-C. (2017). brms: An R package for Bayesian multilevel models using Stan. *Journal of Statistical Software*, *80*(1), 1–28.

Byrd, R. H., Lu, P., Nocedal, J., & Zhu, C. (1995). A limited memory algorithm for bound constrained optimization. *SIAM Journal on scientific computing*, *16*(5), 1190–1208.

Cao, A., Geiger, A., Kreiss, E., Icard, T., & Gerstenberg, T. (2023). A semantics for causing, enabling, and preventing verbs using structural causal models.

Chang, W. (2009). Connecting counterfactual and physical causation. In *Proceedings of the 31th annual conference of the cognitive science society* (pp. 1983–1987). Cognitive Science Society, Austin, TX.

Cheng, P. W., & Novick, L. R. (1991). Causes versus enabling conditions. *Cognition*, *40*, 83–120.

Comrie, B. (1976). The syntax of causative constructions: cross-language similarities and divergences. In *The grammar of causative constructions* (pp. 259–312). Brill.

Cruse, D. A. (1972). A note on english causatives. *Linguistic inquiry*, *3*(4), 520–528.

Degen, J. (2015). Investigating the distribution of some (but not all) implicatures using corpora and web-based methods. *Semantics and Pragmatics*, *8*(11), 1–55.

Degen, J. (2023). The rational speech act framework. *Annual Review of Linguistics*, *9*, 519–540.

De Leeuw, J. R. (2015). jspsych: A javascript library for creating behavioral experiments in a web browser. *Behavior research methods*, *47*, 1–12.

Dowe, P. (2000). *Physical causation.* Cambridge, England: Cambridge University Press.

Fair, D. (1979). Causation and the flow of energy. *Erkenntnis*, *14*(3), 219–250.

Flanigan, J. (2014). A defense of compulsory vaccination. In *Hec forum* (Vol. 26, pp. 5–25).

Fodor, J. A. (1970). Three reasons for not deriving "kill" from "cause to die". *Linguistic Inquiry*, *1*(4), 429–438.

Foot, P. (1967). The problem of abortion and the doctrine of the double effect. *Oxford review*, *5*, 5–15.

Frank, M. C., & Goodman, N. D. (2012). Predicting pragmatic reasoning in language games. *Science*, *336*(6084), 998–998.

Garvey, C., & Caramazza, A. (1974). Implicit causality in verbs. *Linguistic inquiry*, *5*(3), 459–464.

Gerstenberg, T. (2022). What would have happened? counterfactuals, hypotheticals and causal judgements. *Philosophical Transactions of the Royal Society B*, *377*(1866), 20210339.

Gerstenberg, T., Goodman, N. D., Lagnado, D. A., & Tenenbaum, J. B. (2021). A counterfactual simulation model of causal judgments for physical events. *Psychological Review*, *128*(5), 936–975.

Gerstenberg, T., Peterson, M. F., Goodman, N. D., Lagnado, D. A., & Tenenbaum, J. B. (2017). Eye-tracking causality. *Psychological Science*, *28*(12), 1731–1744.

Gerstenberg, T., & Stephan, S. (2021). A counterfactual simulation model of causation by omission. *Cognition*, *216*, 104842.

Gerstenberg, T., & Tenenbaum, J. B. (2017). Intuitive theories. In M. Waldmannn (Ed.), *Oxford handbook of causal reasoning* (pp. 515–548). Oxford University Press.

Gerstenberg, T., Ullman, T. D., Nagel, J., Kleiman-Weiner, M., Lagnado, D. A., & Tenenbaum, J. B. (2018). Lucky or clever? From expectations to responsibility judgments. *Cognition*, *177*, 122-141.

Goldvarg, E., & Johnson-Laird, P. N. (2001). Naive causality: A mental model theory of causal meaning and reasoning. *Cognitive Science*, *25*(4), 565–610.

Goodman, N. D., & Frank, M. C. (2016). Pragmatic language interpretation as probabilistic inference. *Trends in Cognitive Sciences*, *20*(11), 818–829.

Grice, H. P. (1975). Logic and conversation. In P. Cole & J. L. Morgan (Eds.), *Syntax and semantics 3: Speech acts.* New York: Wiley.

Gureckis, T. M., Martin, J., McDonnell, J., Rich, A. S., Markant, D., Coenen, A., . . . Chan, P. (2016). psiturk: An open-source framework for conducting replicable behavioral experiments online. *Behavior research methods*, *48*(3), 829–842.

Hall, N. (2004). Two concepts of causation. In J. Collins, N. Hall, & L. A. Paul (Eds.), *Causation and counterfactuals.* MIT Press.

Hall, N. (2007). Structural equations and causation. *Philosophical Studies*, *132*, 109–136.

Halpern, J. Y. (2016). *Actual causality.* MIT Press.

Halpern, J. Y., & Pearl, J. (2005). Causes and explanations: A structural-model approach. Part I: Causes. *The British Journal for the Philosophy of Science*, *56*(4), 843–887.

Hartshorne, J. K. (2013). What is implicit causality? *Language, Cognition and Neuroscience*, *29*(7), 804–824.

Haspelmath, M. (2016). Universals of causative and anticausative verb formation and the spontaneity scale. *Lingua Posnaniensis*, *58*(2), 33–63.

Henne, P., & O'Neill, K. (2022). Double prevention, causal judgments, and counterfactuals. *Cognitive Science*, *46*(5), e13127.

Hilton, D. J. (1990). Conversational processes and causal explanation. *Psychological Bulletin*, *107*(1), 65–81.

Hirschberg, J. B. (1985). *A theory of scalar implicature (natural languages, pragmatics, inference).* (Doctoral dissertation, University of Pennsylvania)

Hitchcock, C. (1995). Salmon on explanatory relevance. *Philosophy of Science*, *62*(2), 304–320.

Hitchcock, C. (2009). Structural equations and causation: six counterexamples. *Philosophical Studies*, *144*(3), 391–401.

Hitchcock, C., & Knobe, J. (2009). Cause and norm. *Journal of Philosophy*, *11*, 587–612.

Ho, M. K., Abel, D., Correa, C. G., Littman, M. L., Cohen, J. D., & Griffiths, T. L.

(2022). People construct simplified mental representations to plan. *Nature*, *606*(7912), 129–136.

Hobbs, J. R. (2005). Toward a useful concept of causality for lexical semantics. *Journal of Semantics*, *22*(2), 181–209.

Icard, T. F., Kominsky, J. F., & Knobe, J. (2017). Normality and actual causal strength. *Cognition*, *161*, 80–93.

Jackson, C. (2013). Reasonable persons, reasonable circumstances. *San Diego L. Rev.*, *50*, 651–706.

Jara-Ettinger, J., Gweon, H., Schulz, L. E., & Tenenbaum, J. B. (2016). The naïve utility calculus: Computational principles underlying commonsense psychology. *Trends in Cognitive Sciences*, *20*(10), 589–604.

Johnson-Laird, P. N. (1989). *Mental models.* The MIT Press.

Katz, J. J. (1970). Interpretative semantics vs. generative semantics. *Foundations of language*, 220–259.

Kaufmann, S. (2013). Causal premise semantics. *Cognitive Science*, *37*(6), 1136–1170.

Khemlani, S. S., Barbey, A. K., & Johnson-Laird, P. N. (2014). Causal reasoning with mental models. *Frontiers in Human Neuroscience*, *8*.

Kirfel, L., Icard, T., & Gerstenberg, T. (2022). Inference from explanation. *Journal of Experimental Psychology: General*, *151*(7), 1481–1501.

Klettke, B., & Wolff, P. (2003). Differences in how english and german speakers talk and reason about cause. In *Proceedings of the annual meeting of the cognitive science society* (Vol. 25).

Kubricht, J. R., Holyoak, K. J., & Lu, H. (2017). Intuitive physics: Current research and controversies. *Trends in Cognitive Sciences*, *21*(10), 749–759.

Lagnado, D. A., Fenton, N., & Neil, M. (2013). Legal idioms: a framework for evidential reasoning. *Argument & Computation*, *4*(1), 46–63.

Lagnado, D. A., Gerstenberg, T., & Zultan, R. (2013). Causal responsibility and

counterfactuals. *Cognitive Science*, *47*, 1036–1073.

Langenhoff, A. F., Wiegmann, A., Halpern, J. Y., Tenenbaum, J. B., & Gerstenberg, T. (2021). Predicting responsibility judgments from dispositional inferences and causal attributions. *Cognitive Psychology*, *129*, 101412.

Levin, B., & Hovav, M. R. (1994). A preliminary analysis of causative verbs in english. *Lingua*, *92*, 35–77.

Lewis, D. (1973). Causation. *The Journal of Philosophy*, *70*(17), 556–567.

Lombrozo, T. (2010). Causal-explanatory pluralism: How intentions, functions, and mechanisms influence causal ascriptions. *Cognitive Psychology*, *61*(4), 303–332.

Machamer, P., Darden, L., & Craver, C. F. (2000). Thinking about mechanisms. *Philosophy of Science*, *67*(1), 1–25.

Mackie, J. L. (1974). *The cement of the universe*. Oxford: Clarendon Press.

Malm, H. M. (1989). Killing, letting die, and simple conflicts. *Philosophy & public affairs*, 238–258.

Matsumoto, Y. (1997). Scales, implicatures, and in fact, if not, and let alone constructions. *Studies in English*, 685–699.

Mayol, L., & Castroviejo, E. (2013). How to cancel an implicature. *Journal of Pragmatics*, *50*(1), 84–104.

Mayrhofer, R., & Waldmann, M. R. (2016). Causal agency and the perception of force. *Psychonomic Bulletin & Review*, *23*(3), 789–796.

McCawley, J. D. (1978). Conversational implicature and the lexicon. *Syntax and semantics*, *9*, 245–259.

McDermott, M. (1995). Redundant causation. *British Journal for the Philosophy of Science*, *46*, 523–544.

McGrath, S. (2003). Causation and the making/allowing distinction. *Philosophical Studies*, *114*(1), 81–106.

McMahan, J. (1993). Killing, letting die, and withdrawing aid. *Ethics*, 250–279.

Nadathur, P., & Lauer, S. (2020). Causal necessity, causal sufficiency, and the implications of causative verbs. *Glossa: a journal of general linguistics*, *5*(1), 1–37.

Niemi, L., Hartshorne, J., Gerstenberg, T., Stanley, M., & Young, L. (2020). Moral values reveal the causality implicit in verb meaning. *Cognitive Science*, *44*(6), e12838.

Pearl, J. (2000). *Causality: Models, reasoning and inference.* Cambridge, England: Cambridge University Press.

Potts, C., Lassiter, D., Levy, R., & Frank, M. C. (2015). Embedded implicatures as pragmatic inferences under compositional lexical uncertainty. *Journal of Semantics*, *33*(4), 755–802.

R Core Team. (2019). R: A language and environment for statistical computing [Computer software manual]. Vienna, Austria.

Rodríguez-Arias, D., Rodriguez Lopez, B., Monasterio-Astobiza, A., & Hannikainen, I. R. (2020). How do people use 'killing','letting die'and related bioethical concepts? contrasting descriptive and normative hypotheses. *Bioethics*, *34*(5), 509–518.

Rose, D., Sievers, E., & Nichols, S. (2021). Cause and burn. *Cognition*, *207*(104517).

Rozenblit, L., & Keil, F. (2002). The misunderstood limits of folk science: An illusion of explanatory depth. *Cognitive science*, *26*(5), 521–562.

Salmon, W. C. (1984). *Scientific explanation and the causal structure of the world.* Princeton University Press, Princeton NJ.

Schaffer, J. (2000). Causation by disconnection. *Philosophy of Science*, *67*(2), 285.

Schaffer, J. (2013). Causal contextualisms. In *Contrastivism in philosophy* (pp. 43–71). Routledge.

Shibatani, M. (1976). The grammar of causative constructions: A conspectus. In *The grammar of causative constructions* (pp. 1–40). Brill.

Shin, S. M., & Gerstenberg, T. (2023). Learning what matters: Causal abstraction in human inference. In M. B. Goldwater, F. Anggoro, B. Hayes, & D. C. Ong (Eds.), *Proceedings of the 45th Annual Conference of the Cognitive Science Society.*

Sloman, S. A. (2005). *Causal models: How people think about the world and its alternatives.* Oxford University Press, USA.

Sloman, S. A., Barbey, A. K., & Hotaling, J. M. (2009). A causal model theory of the meaning of cause, enable, and prevent. *Cognitive Science*, *33*(1), 21–50.

Smith, C. S. (1970). Jespersen's 'move and change' class and causative verbs in english. *Linguistic and literary studies in honor of Archibald A. Hill*, *2*, 101–109.

Smith, K. A., Hamrick, J. B., Sanborn, A. N., Battaglia, P. W., Gerstenberg, T., Ullman, T. D., & Tenenbaum, J. B. (in press). Probabilistic models of physical reasoning. In T. L. Griffiths, N. Chater, & J. B. Tenenbaum (Eds.), *Reverse engineering the mind: Probabilistic models of cognition.*

Smith, K. A., & Vul, E. (2013). Sources of uncertainty in intuitive physics. *Topics in Cognitive Science*, *5*(1), 185–199.

Sosa, F. A., Ullman, T., Tenenbaum, J. B., Gershman, S. J., & Gerstenberg, T. (2021). Moral dynamics: Grounding moral judgment in intuitive physics and intuitive psychology. *Cognition*, *217*, 104890.

Summers, A. (2018). Common-sense causation in the law. *Oxford Journal of Legal Studies*, *38*(4), 793–821.

Talmy, L. (1988). Force dynamics in language and cognition. *Cognitive Science*, *12*(1), 49–100.

Tavares, Z., Koppel, J., Zhang, X., Das, R., & Solar-Lezama, A. (2021, 18–24 Jul). A language for counterfactual generative models. In M. Meila & T. Zhang (Eds.), *Proceedings of the 38th international conference on machine learning* (Vol. 139, pp. 10173–10182). PMLR.

Thomson, J. J. (1976a). A defense of abortion. In *Biomedical ethics and the law* (pp. 39–54). Springer.

Thomson, J. J. (1976b). Killing, letting die, and the trolley problem. *The Monist*, *59*(2), 204–217.

Tobia, K. P. (2018). How people judge what is reasonable. *Alabama Law Review*, *70*, 293–359.

Ullman, T. D., Spelke, E., Battaglia, P., & Tenenbaum, J. B. (2017). Mind games: Game engines as an architecture for intuitive physics. *Trends in Cognitive Sciences*, *21*(9), 649–665.

Ullman, T. D., Tenenbaum, J. B., Baker, C. L., Macindoe, O., Evans, O. R., & Goodman, N. D. (2009). Help or hinder: Bayesian models of social goal inference. In *Advances in Neural Information Processing Systems* (Vol. 22, pp. 1874–1882).

Virtanen, P., Gommers, R., Oliphant, T. E., Haberland, M., Reddy, T., Cournapeau, D., . . . SciPy 1.0 Contributors (2020). SciPy 1.0: Fundamental Algorithms for Scientific Computing in Python. *Nature Methods*, *17*, 261–272.

White, P. A. (2014). Singular clues to causality and their use in human causal judgment. *Cognitive Science*, *38*(1), 38–75.

Wierzbicka, A. (1975). Why "kill" does not mean "cause to die": the semantics of action sentences. *Foundations of language*, *13*(4), 491–528.

Wolff, P. (2003). Direct causation in the linguistic coding and individuation of causal events. *Cognition*, *88*(1), 1–48.

Wolff, P. (2007). Representing causation. *Journal of Experimental Psychology: General*, *136*(1), 82–111.

Wolff, P., Barbey, A. K., & Hausknecht, M. (2010). For want of a nail: How absences cause events. *Journal of Experimental Psychology: General*, *139*(2), 191–221.

Wolff, P., Jeon, G.-H., & Li, Y. (2009). Causers in english, korean, and chinese and the individuation of events. *Language and Cognition*, *1*(2), 167–196.

Wolff, P., Klettke, B., Ventura, T., & Song, G. (2005). Expressing causation in english and other languages. In W. Ahn, R. L. Goldstone, B. C. Love, A. B. Markman, & P. Wolff (Eds.), *Categorization inside and outside the laboratory: Essays in honor of Douglas L. Medin* (p. 29-48). American Psychological Association.

Wolff, P., & Song, G. (2003). Models of causation and the semantics of causal verbs. *Cognitive Psychology*, *47*(3), 276–332.

Woodward, J. (2003). *Making things happen: A theory of causal explanation.* Oxford, England: Oxford University Press.

Woodward, J. (2006). Sensitive and insensitive causation. *The Philosophical Review*, *115*(1), 1–50.

Woodward, J. (2021). *Causation with a human face: Normative theory and descriptive psychology.* Oxford University Press.

Wu, S., Sridhar, S., & Gerstenberg, T. (2022). That was close! a counterfactual simulation model of causal judgments about decisions. In *Proceedings of the annual meeting of the cognitive science society* (Vol. 44).

Wu, S. A., & Gerstenberg, T. (2024). If not me, then who? responsibility and replacement. *Cognition*, *242*, 105646.

Wu, S. A., Sridhar, S., & Gerstenberg, T. (2023). A computational model of responsibility judgments from counterfactual simulations and intention inferences. In M. B. Goldwater, F. Anggoro, B. Hayes, & D. C. Ong (Eds.), *Proceedings of the 45th Annual Conference of the Cognitive Science Society.*

Zhou, L., Smith, K. A., Tenenbaum, J. B., & Gerstenberg, T. (2023). Mental Jenga: A counterfactual simulation model of causal judgments about physical support. *Journal of Experimental Psychology: General*.

Appendix A

Sentence frames from norming study (Experiment 1A.
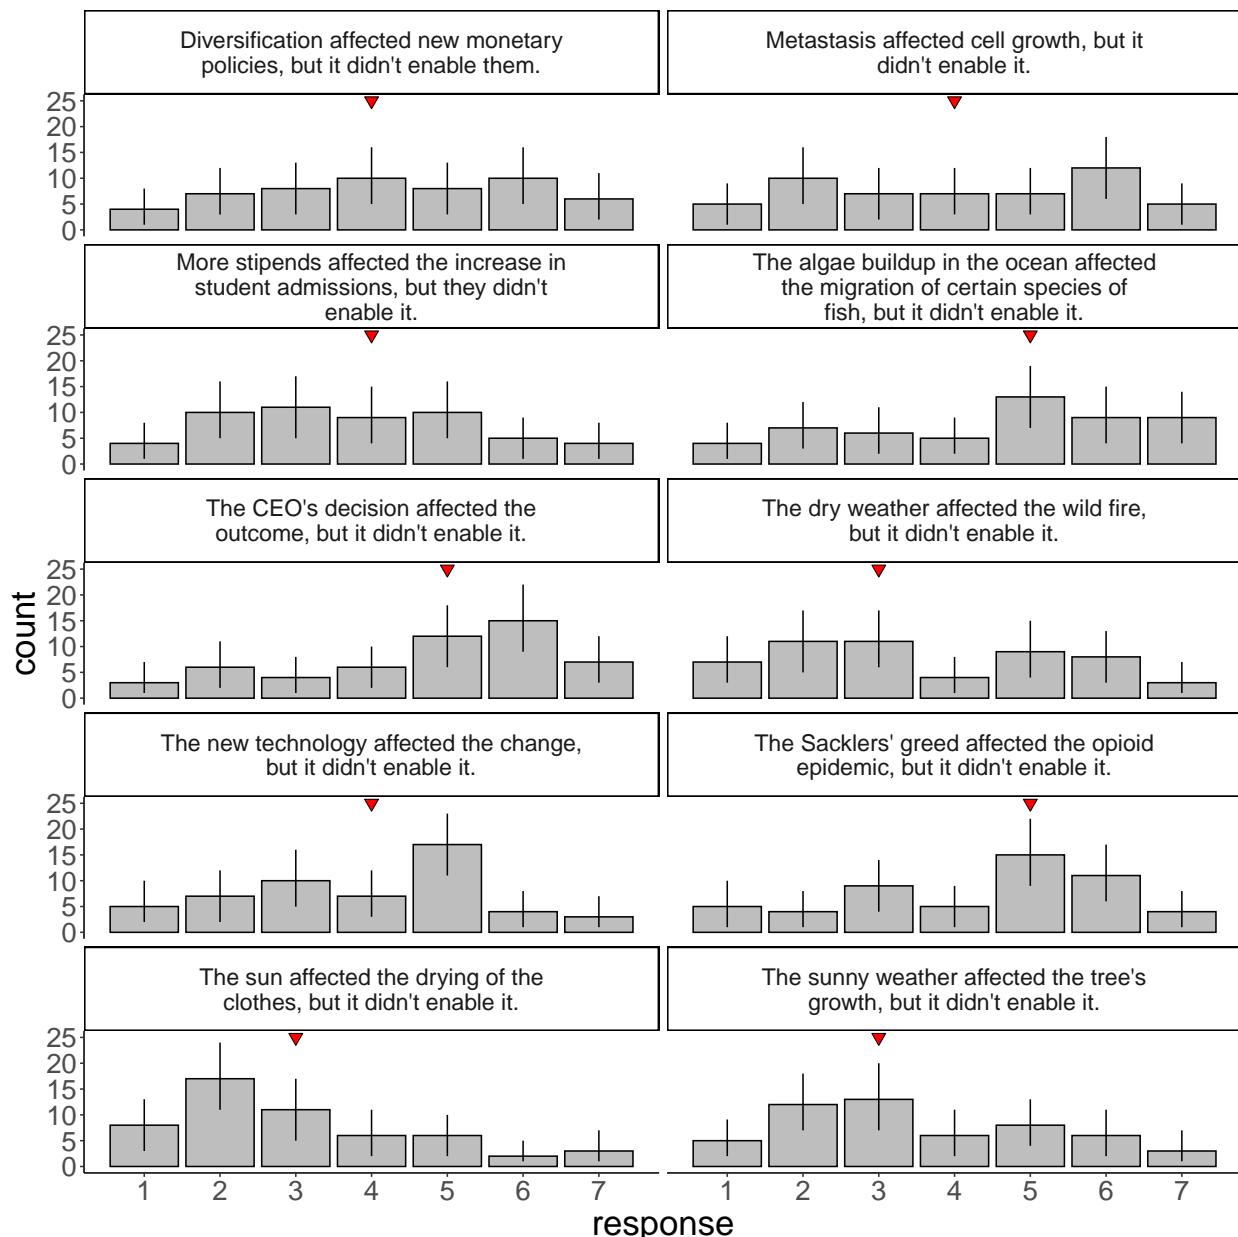
Table A1
All sentence frames from norming study. Included sentences had median ratings for all verbs above the midpoint of the scale. Excluded sentences had median ratings at the midpoint or below for at least one causal expression.

Included Sentences
1. The dry weather _____ the wild fire.
2. The CEO's decision _____ the outcome.
3. The new technology _____ the change.
4. The Sackler's greed _____ the opioid epidemic.
5. The sunny weather _____ the tree's growth.
6. More stipends _____ the increase in student admissions.
7. The sun _____ the drying of the clothes.
8. Metastasis _____ cell growth.
9. Diversification _____ new monetary policies.
10. The algae buildup in the ocean _____ the migration of certain species of fish.

Excluded Sentences
11. The collapse of Lehman Brothers _____ the financial crisis.
12. The zoning restrictions _____ the housing shortage.
13. The breaking of the dam _____ the flood.
14. Janelle's working hard _____ her success.
15. The new traffic signs _____ the decrease in fatalities.
16. The striker deflecting the ball _____ the goal.
17. Turning off life support _____ the patient's death.
18. Deforestation _____ wildlife displacement.
19. Erosion _____ density loss.
20. The construction at the intersection _____ the traffic in the vicinity.

Attention Checks
1. The lightening strike caused the fire.
2. The crank caused him to open the window.
3. Receiving the loan enabled her to buy the house.
4. The cold breeze enabled him to close the window.
5. The distracting noise affected his performance.
6. The earthquake affected the building to fall.

Appendix B
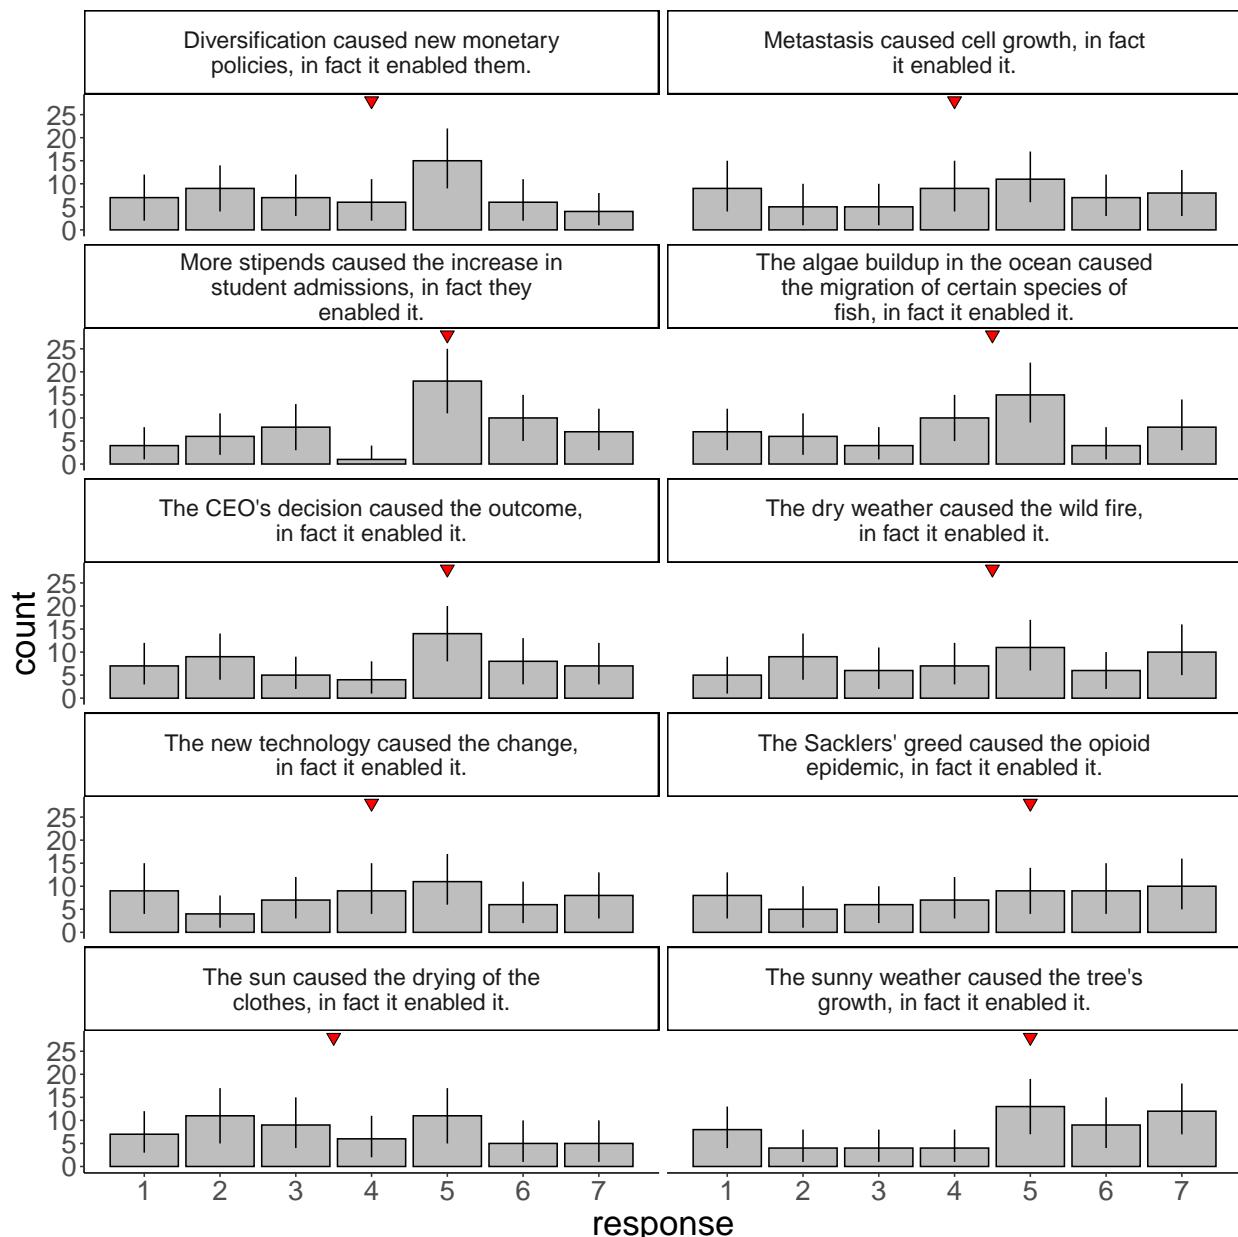
Experiment 1B Affected → Enabled Items

Figure B1. Histograms of participant responses in Experiment 1B for all items with the affected → enabled order. *Note*: 1 = "definitely not acceptable", 4 = "unsure", 7 = "definitely acceptable". There is substantial variety in the distributions of participants' responses, with some frames skewing toward the acceptable side of the scale and others skewing to the unacceptable side. Median rating is indicated by red triangle. Error bars represent bootstrapped 95% confidence intervals.
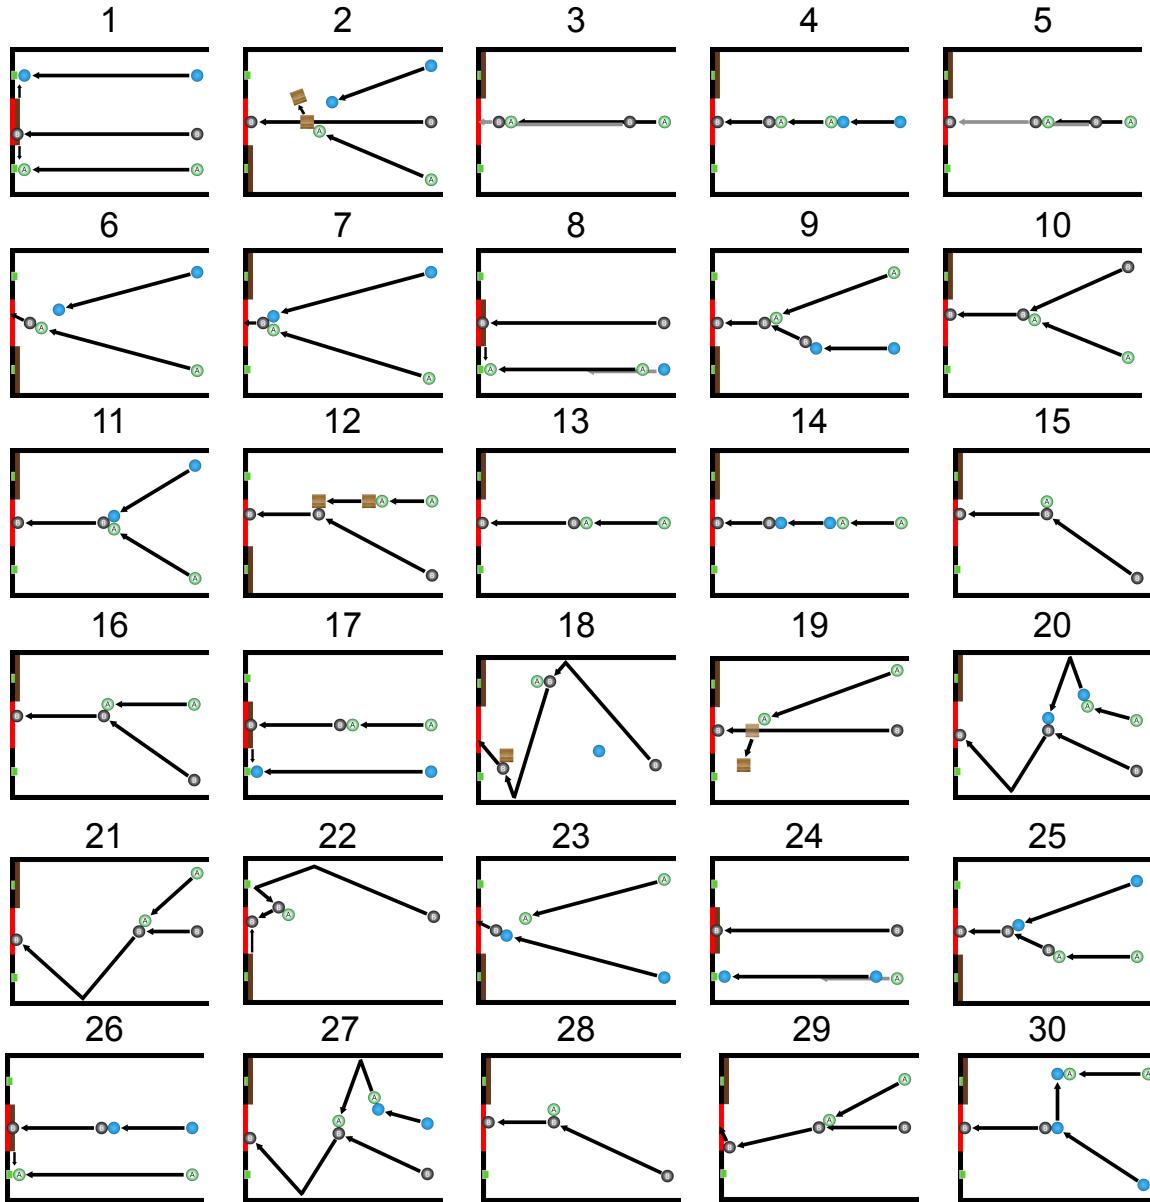
Appendix C

Experiment 1C Caused → Enabled Items

Figure C1. Histograms of participant responses in Experiment 1C for all items with the caused → enabled order. *Note*: 1 = "definitely not acceptable", 4 = "unsure", 7 = "definitely acceptable". Again we see substantial variability within and between frames. Some response distributions appear more uniform, while others skew more to acceptability. Median rating is indicated by red triangle. Error bars represent bootstrapped 95% confidence intervals.

Appendix D

Scenario Schematics and Aspect Values


Figure D1. Trial schematics for all video scenarios in Experiment 2 and Experiment 3. Video clips are available here: `https://github.com/cicl-stanford/causal_language/tree/master/code/experiments/experiment2/static/videos/mp4`

Table D1

Table of the aspect values for the scenarios from Experiment 2 and 3. Scenario numbers correspond to the numbers in Figure D1. Aspects are computed with noise value $\theta = 1.0$; the optimal value found for the full model in Experiment 2.

Scenario	Whether	How	Sufficient	Moving	Unique
1	0.0	0	0.0	1	0
2	0.299	0	1.0	1	0
3	0.0	1	0.0	1	1
4	0.399	1	0.0	1	1
5	0.063	1	0.06	1	1
6	0.0	1	1.0	1	1
7	0.0	1	1.0	1	0
8	0.111	0	1.0	1	0
9	0.583	1	0.0	1	0
10	0.762	1	0.759	1	1
11	0.88	1	0.059	1	0
12	0.997	1	0.003	1	0
13	1.0	1	1.0	1	1
14	1.0	1	1.0	1	0
15	0.999	1	0.999	0	1
16	0.998	1	0.995	1	1
17	1.0	1	0.0	1	1
18	0.311	1	0.05	0	0
19	1.0	0	1.0	1	0
20	0.951	1	0.047	1	0
21	0.289	1	0.277	1	1
22	1.0	1	1.0	0	1
23	0.0	0	0.0	1	0
24	0.0	0	0.0	1	0
25	1.0	1	0.396	1	0
26	1.0	0	0.0	1	0
27	0.955	1	0.044	1	1
28	0.8	1	0.769	0	1
29	0.092	1	0.098	1	1
30	0.0	0	0.0	1	0

Appendix E

Experiment 2: Trial Response Distributions

Table E1

Distribution of participant responses for all trials in Experiment 2. Scenario numbers correspond to schematics in Figure D1. Note: Some rows do not sum to 1 due to rounding.

Scenario	No Difference	Affected	Enabled	Caused
1	0.95	0.02	0.02	0.02
2	0.10	0.11	0.76	0.03
3	0.63	0.24	0.03	0.10
4	0.05	0.35	0.11	0.48
5	0.58	0.29	0.10	0.03
6	0.02	0.06	0.10	0.82
7	0.06	0.44	0.19	0.31
8	0.21	0.05	0.73	0.02
9	0.03	0.31	0.26	0.40
10	0.08	0.16	0.11	0.65
11	0.03	0.48	0.19	0.29
12	0.02	0.26	0.53	0.19
13	0.00	0.02	0.06	0.92
14	0.02	0.21	0.19	0.58
15	0.02	0.37	0.37	0.24
16	0.00	0.19	0.15	0.66
17	0.02	0.13	0.11	0.74
18	0.10	0.66	0.18	0.06
19	0.03	0.03	0.89	0.05
20	0.10	0.35	0.29	0.26
21	0.18	0.58	0.06	0.18
22	0.03	0.24	0.48	0.24
23	0.98	0.00	0.02	0.00
24	0.97	0.00	0.02	0.02
25	0.03	0.50	0.29	0.18
26	0.16	0.05	0.77	0.02
27	0.00	0.47	0.11	0.42
28	0.08	0.35	0.35	0.21
29	0.39	0.50	0.00	0.11
30	0.98	0.00	0.00	0.02

Appendix F

No Semantics and No Pragmatics: Top Regression Fit

Table F1
Fixed effects of the top performing No Semantics and No Pragmatics Model. The Estimate column gives the posterior mean for the given term, while the Estimate Error gives the posterior variance. The columns CI Lower Bound and CI Upper Bound give the lower and upper bounds of the 95% credible interval for the given term.

Term	Estimate	Estimate Error	CI Lower Bound	CI Upper Bound
Threshold No Difference|Affected	2.42	0.64	1.16	3.67
Threshold Affected|Enabled	3.74	0.64	2.47	4.98
Threshold Enabled|Caused	4.76	0.64	3.48	6.01
whether	1.58	0.42	0.79	2.43
how	1.46	0.46	0.58	2.39
sufficient	1.54	0.37	0.82	2.29
moving	1.12	0.48	0.16	2.07
unique	0.16	0.39	-0.59	0.93

Appendix G

Semantics Analysis Summary

Table G1

Summary of our semantics evaluation. We considered all the possible ways we could assign definitions to "caused", "enabled", and "affected", and then restricted this set to just those assignments that were consistent with our findings experiment 1. This left us with a total of 258 definition assignments where "caused" implied "enabled" and "enabled" in turn implied "affected". We then evaluated our full model with each of these definition assignments, fixing θ based on prior work and re-fitting ν and λ for each assignment. The top five and bottom five definition assignments are listed in the table, with the log likelihood of the data under the given model. Notably, our specified model semantics is the top performing assignment for our data.

Rank	Caused	Enabled	Affected	Log Likelihood
1	$(\mathcal{W} \vee \mathcal{S}) \wedge \mathcal{H}$	$\mathcal{W} \vee \mathcal{S}$	$\mathcal{W} \vee \mathcal{S} \vee \mathcal{H}$	-1897.71
2	$(\mathcal{W} \vee \mathcal{S}) \wedge \mathcal{H}$	$(\mathcal{W} \wedge \mathcal{H}) \vee \mathcal{S}$	$\mathcal{H} \vee \mathcal{S}$	-1919.90
3	$(\mathcal{W} \vee \mathcal{S}) \wedge \mathcal{H}$	$(\mathcal{H} \wedge \mathcal{S}) \vee \mathcal{W}$	$\mathcal{W} \vee \mathcal{H}$	-1921.74
4	$(\mathcal{H} \wedge \mathcal{S}) \vee \mathcal{W}$	$\mathcal{W} \vee \mathcal{S}$	$\mathcal{W} \vee \mathcal{S} \vee \mathcal{H}$	-1956.87
5	$(\mathcal{W} \wedge \mathcal{H}) \vee \mathcal{S}$	$\mathcal{W} \vee \mathcal{S}$	$\mathcal{W} \vee \mathcal{H} \vee \mathcal{S}$	-1978.29
		$\ldots$		
254	$\mathcal{W} \wedge \mathcal{S}$	$(\mathcal{W} \vee \mathcal{H}) \wedge \mathcal{S}$	$\mathcal{S}$	-2291.64
255	$\mathcal{W} \wedge \mathcal{H} \wedge \mathcal{S}$	$\mathcal{H} \wedge \mathcal{S}$	$(\mathcal{W} \wedge \mathcal{H}) \vee \mathcal{S}$	-2294.30
256	$\mathcal{W} \wedge \mathcal{H} \wedge \mathcal{S}$	$\mathcal{H} \wedge \mathcal{S}$	$(\mathcal{W} \vee \mathcal{H}) \wedge \mathcal{S}$	-2296.99
257	$\mathcal{W} \wedge \mathcal{H} \wedge \mathcal{S}$	$\mathcal{H} \wedge \mathcal{S}$	$\mathcal{W} \vee \mathcal{S}$	-2303.57
258	$\mathcal{W} \wedge \mathcal{H} \wedge \mathcal{S}$	$\mathcal{H} \wedge \mathcal{S}$	$\mathcal{S}$	-2332.10

Appendix H

Experiment 3: Trial Pairings and Participant Summary Responses

Table H1

Video pairings for all trials in Experiment 3. The headers indicate trial expression. Video numbers refer to the scenarios in Figure D1. Means less than 0.5 indicate overall ratings favoring video 1, while means greater than 0.5 indicate overall ratings favoring video 2.

Trial	Video 1	Video 2	Mean	SD
Caused				
1	10	17	0.50	0.30
2	4	9	0.55	0.28
3	10	16	0.54	0.23
4	2	14	0.67	0.30
5	10	20	0.35	0.29
6	15	16	0.68	0.24
7	6	30	0.13	0.23
8	13	19	0.19	0.22
9	13	25	0.22	0.26
Enabled				
10	2	19	0.58	0.21
11	15	22	0.50	0.24
12	12	22	0.44	0.27
13	1	22	0.79	0.25
14	12	29	0.25	0.27
15	11	26	0.7 0	0.29
16	13	19	0.78	0.27
17	8	24	0.17	0.24
18	4	26	0.73	0.28
Affected				
19	17	28	0.69	0.28
20	15	16	0.44	0.24
21	5	7	0.69	0.25
22	14	21	0.48	0.28
23	25	26	0.42	0.30
24	7	23	0.21	0.27
25	18	24	0.21	0.27
26	18	23	0.19	0.23
27	18	30	0.14	0.20
Made No Difference				
28	3	29	0.28	0.27
29	23	30	0.49	0.23
30	5	29	0.29	0.25
31	3	12	0.20	0.29
32	3	14	0.20	0.29
33	5	7	0.28	0.28
34	8	24	0.72	0.31
35	1	21	0.23	0.29
36	1	27	0.18	0.28